

Servers & Security Groups

July 6, 2020

1 Security Groups

The following is the syntax required to create a `SecurityGroup`:

Security groups are associated with specific resources and not tied to the subnets

Type: `AWS::EC2::SecurityGroup`

Properties:

 GroupDescription: String

 GroupName: String

 SecurityGroupEgress:

 - Egress

 SecurityGroupIngress:

 - Ingress

 Tags:

 - Tag

 VpcId: String

Although they are not required, the `SecurityGroupEgress` and `SecurityGroupIngress` property rules are the most critical to the `SecurityGroup` as it defines where the traffic will go. While `SecurityGroupEgress` defines outbound traffic, `SecurityGroupIngress` defines the inbound traffic

1.1 Ingress Rules and Egress Rules

- Ingress rules are for inbound traffic, and egress rules are for the outbound traffic
- Ingress rules restrict or allow traffic trying to reach our resources on specific ports
- Egress rules restrict or allow traffic originating from our server –typically we are ok allowing all outbound traffic without restrictions as this doesn't pose a risk for a security breach

1.2 Traffic is blocked by default

- In cloud traffic is completely blocked, so you have to explicitly open ports to allow traffic in and out. This is a general networking concept

1.3 Limit inbound traffic for security

- For ingress rules, we want to limit inbound traffic, for security, to a single port or just a handful of ports required by the application we are running
- If it's a public web server, for example, it will require `port 80` open to the world (World = `0.0.0.0/0`)

- Should you need the SSH port open, restrict this port only to specify IP address
- should you need the SSH port open, restrict this port only to your specific IP address

1.4 For outbound traffic, give full access

For egress rules, we want to give this resource full access to the internet so we give egress access to all ports, from 0 all the way to 65535

2 Creating AutoScaling Group

The autoscaling group must be given a launch configuration and a scaling policy

The following is the syntax used for AutoScaling LaunchConfiguration:

Type: `AWS::AutoScaling::LaunchConfiguration`

Properties:

```
AssociatePublicIpAddress: Boolean
BlockDeviceMappings:
  - BlockDeviceMapping
ClassicLinkVPCId: String
ClassicLinkVPCSecurityGroups:
  - String
EbsOptimized: Boolean
IamInstanceProfile: String
ImageId: String
InstanceId: String
InstanceMonitoring: Boolean
InstanceType: String
KernelId: String
KeyName: String
LaunchConfigurationName: String
PlacementTenancy: String
RamDiskId: String
SecurityGroups:
  - String
SpotPrice: String
UserData: String
```

The `ImageId` and `Instance Type` are the only required properties for a `LaunchConfiguration`. However, there are many useful properties you will likely want to include

In the example below we have done the following: - Specified 10gbs for our `VolumeSize` - Referenced the previously defined `WebServerSecGroup` for our `SecurityGroup` - Set our `InstanceType` to `t3.medium` for our EC2 Instance

WebAppLaunchConfig:

```
Type: AWS::AutoScaling::LaunchConfiguration
Properties:
  UserData:
    Fn::Base64: !Sub |
```

```

#!/bin/bash
apt-get update -y
apt-get install unzip awscli -y
apt-get install apache2 -y
systemctl start apache2.service
cd /var/www/html
aws s3 cp s3://udacity-demo-1/udacity.zip .
unzip -o udacity.zip
ImageId: ami-005bdb005fb00e791
IamInstanceProfile: !Ref ProfileWithRolesForOurApp
SecurityGroups:
- Ref: WebServerSecGroup
InstanceType: t3.small
BlockDeviceMappings:
- DeviceName: "/dev/sdk"
Ebs:
  VolumeSize: '10'

```

The following is the required syntax for TargetGroup

Type: AWS::ElasticLoadBalancingV2::TargetGroup

Properties:

```

HealthCheckEnabled: Boolean
HealthCheckIntervalSeconds: Integer
HealthCheckPath: String
HealthCheckPort: String
HealthCheckProtocol: String
HealthCheckTimeoutSeconds: Integer
HealthyThresholdCount: Integer
Matcher:
  Matcher
Name: String
Port: Integer
Protocol: String
Tags:
- Tag
TargetGroupAttributes:
- TargetGroupAttribute
TargetType: String
Targets:
- TargetDescription
UnhealthyThresholdCount: Integer
VpcId: String

```

Health Checks are the requests your Application Load Balancer sends to its registered targets. These periodic requests test the status of these targets. You can see us defining our Health Check properties in the example below:

WebAppTargetGroup:

Type: AWS::ElasticLoadBalancingV2::TargetGroup

```

Properties:
  HealthCheckIntervalSeconds: 35
  HealthCheckPath: /
  HealthCheckProtocol: HTTP
  HealthCheckTimeoutSeconds: 30
  HealthyThresholdCount: 2
  Port: 80
  Protocol: HTTP
  UnhealthyThresholdCount: 5
  VpcId:
    Fn::ImportValue:
      Fn::Sub: "${EnvironmentName}-VPCID"

```

In the above example we specify the following: - The port where our targets receive traffic - Port: 80 - The protocol the load balancer uses when performing health checks on targets - HealthCheckProtocol: HTTP - The time it takes to determine a non-responsive target is unhealthy-HealthCheckingIntervalSeconds: 35 - The number of health/unhealthy checks required to change the health status - HealthyThresholdCount: 2 UnhealthyThreshholdCount: 5

3 Security Group Error

For the load balancer, we basically have port 80 as both inbound and outbound

Thus your basically limiting the load balancer to only be able to talk with HTTP

That is why we add port 8080 as shown below

| Type | Protocol | Port Range | Destination | Description |
|------------|----------|------------|------------------|----------------------------|
| HTTP | TCP | 80 | Custom 0.0.0.0/0 | e.g. SSH for Admin Desktop |
| Custom TCP | TCP | 8080 | Custom 0.0.0.0/0 | e.g. SSH for Admin Desktop |
| Custom TCP | TCP | 8080 | Custom ::/0 | e.g. SSH for Admin Desktop |

Add Rule