

Chapter 05 - Classes and Interfaces

May 8, 2021

0.1 Item 37: Compose Classes Instead of Nesting Many Levels of Built-in Types

- python's built-in dictionary type is wonderful for maintaining dynamic internal state over the lifetime of an object
- by **dynamic** we mean situations where you need to do bookkeeping for unexpected set of identifiers
- imagine you need to record the grades of a set of students whose names aren't known in advance
- **Nesting Many Levels of Built-in Types Means don't put a dictionary in another dictionary or don't put a tuple in a dictionary inside of a dictionary**
- dictionaries and their related built-in types are easy to use that there's a danger of overextending them to write brittle code
- for example, say I want to extend the student/grade class feature to keep a **list** of grades by subject, not just overall
- I can do this by changing the **_grade** dictionary to map student names (its keys) to yet another dictionary (its values)
- the inner most dictionary will map subjects (its keys) to a **list** of grades (its values)
- we will use a **defaultdict** instance for the inner dictionary to handle missing subjects
- imagine that the requirements change again
- we also want to track the weight of each score toward the overall grade in the class so that midterm and final exams are more important than pop quizzes
- instead of mapping subjects (its keys) to a **list** of grades (its values), I can use the tuple of (score, weight) in the values list

```
[8]: from collections import defaultdict

class WeightedGradebook:
    def __init__(self):
        self._grades = {}

    def add_student(self, name):
        self._grades[name] = defaultdict(list)

    def report_grade(self, name, subject, score, weight):
        by_subject = self._grades[name]
        grade_list = by_subject[subject]
        grade_list.append((score, weight))
```

```

def average_grade(self, name):
    by_subject = self._grades[name]
    score_sum, score_count = 0, 0
    for subject, scores in by_subject.items():
        subject_avg, total_weight = 0, 0

        for score, weight in scores:
            subject_avg += score * weight
            total_weight += weight

        score_sum += subject_avg / total_weight
        score_count += 1

    return score_sum / score_count

book = WeightedGradebook()
book.add_student('Albert Einstein')
book.report_grade('Albert Einstein', 'Math', 75, 0.05)
book.report_grade('Albert Einstein', 'Math', 65, 0.15)
book.report_grade('Albert Einstein', 'Math', 70, 0.80)
book.report_grade('Albert Einstein', 'Gym', 100, 0.40)
book.report_grade('Albert Einstein', 'Gym', 85, 0.60)
print(book.average_grade('Albert Einstein'))

```

80.25

- as you can imagine the complexity is just going to get unmanageable
- this is when we need to make the leap from built-in types like dictionaries, tuples, sets and lists to a hierarchy of classes
- when you realize that the class is getting complicated, break it all out into classes
- you can then provide well-defined interfaces that better encapsulate your data
- this lets you create a layer of abstraction between your interfaces and your concrete implementations

0.1.1 Refactoring to Classes

- we can start moving to classes at the bottom of the dependency tree:
 - a single grade
- a class seems too heavyweight for such simple information
- a tuple though, seems appropriate because grades are immutable
- below I use the tuple of (score, weight) to track grades in a list

```

[9]: grades = []
grades.append((95, 0.45))
grades.append((85, 0.55))
total = sum(score * weight for score, weight in grades)
total_weight = sum(weight for _, weight in grades)

```

```
average_grade = total / total_weight
```

- note the `_` is used in loops where we want to ignore the unused variable
- the problem with the code above is that the `tuple` instances are positional
- if we want to associate more information with a grade, such as a set of notes from the teacher, we need to rewrite every uage of the two-tuple to be aware that there are now three items present instead
- this means we now have to use multiple `_` in the code

```
[10]: grades = []
grades.append((95, 0.45, 'Great job'))
grades.append((85, 0.55, 'Better next time'))
total = sum(score * weight for score, weight, _ in grades)
total_weight = sum(weight for _, weight, _ in grades)
average_grade = total / total_weight
```

- the pattern of extending tuples longer and longer is similar to deepening layers of dictionaries
- as soon as you find yourself going longer than a two-tuple, it's time to consider another approach
- as soon as you see yourself going longer than a two tuple, its time to reconsider your approach
- the `namedtuple` type in the `collections` built-in module does exactly what we need in this case
- it lets us easily define tiny, immutable data classes

```
[11]: from collections import namedtuple

Grade = namedtuple('Grade', ('score', 'weight'))
```

- these classes can be constructed with positional or keyword arguments
- the fields are accessible with named attributes
- having named attributes makes it easy to move from a `namedtuple` to a class later if the requirements change again and we need to say, support mutability or behaviors in the simple data containers

Limitations of namedtuple: - you cant specify default argument values for `namedtuple` class - this makes them unwieldy when your data may have many optional properties. If you find yourself using more than a handful of attributes, using the built in `dataclass` module may be a better choice - the attribute values of `namedtuple` instances are still accessible using numerical indexes and iterations - especially in externalized APIs, this can lead to unintentional usage that makes it harfer to mvove to a real `class` later - if your not in control of all of the usages of your `namedtuple` instances, its better to explicitly define a new class

- Next, we write a class to represent a single subject that contains a `set` of grades

```
[12]: class Subject:
    def __init__(self):
        self._grades = []

    def report_grade(self, score, weight):
```

```

        self._grades.append(Grade(score, weight))

    def average_grade(self):
        total, total_weight = 0, 0
        for grade in self._grades:
            total += grade.score * grade.weight
            total_weight += grade.weight
        return total / total_weight

```

- then we write a class to represent a set of subjects that are being studied by a single student

```

[13]: class Student:
    def __init__(self):
        self._subjects = defaultdict(Subject)

    def get_subject(self, name):
        return self._subjects[name]

    def average_grade(self):
        total, count = 0, 0
        for subject in self._subjects.values():
            total += subject.average_grade()
            count += 1
        return total / count

```

- finally we write a container for all the students, keyed dynamically by their names

```

[14]: class Gradebook:
    def __init__(self):
        self._students = defaultdict(Student)

    def get_student(self, name):
        return self._students[name]

```

- the line count of these classes is almost double the previous implementation size, but this code is much easier to read

```

[15]: book = Gradebook()
albert = book.get_student('Albert Einstein')
math = albert.get_subject('Math')
math.report_grade(75, 0.05)
math.report_grade(65, 0.15)
math.report_grade(70, 0.80)
gym = albert.get_subject('Gym')
gym.report_grade(100, 0.40)
gym.report_grade(85, 0.60)
print(albert.average_grade())

```

80.25

- avoid making dictionaries with values that are dictionaries, long tuples or complex nesting of other built-in types
- use `namedtuple` for lightweight immutable data containers before you need the flexibility of a full class
- move your code to using multiple classes when your internal state dictionaries get complicated

0.2 Item 38: Accept Functions Instead of Classes for Simple Interfaces

- many of python's built-in APIs allow you to customize behavior by passing in a function
- these **hooks** are used by APIs to call back your code while they execute
- for example the `list` type's `sort` method takes an optional `key` argument that's used to determine each index's value for sorting
- Here we sort a `list` of names based on their lengths by providing the `len` built-in function as the `key` hook

```
[16]: names = ['Socrates', 'Archimedes', 'Plato', 'Aristotle']
names.sort(key=len)
print(names)
```

```
['Plato', 'Socrates', 'Aristotle', 'Archimedes']
```

- in other languages, you might expect hooks to be defined by an abstract class
- in python, many hooks are just stateless functions with well-defined arguments and return values
- functions are ideal for hooks because they are easier to describe and simpler to define than classes
- functions work as hooks because Python has **first-class** functions: Functions and methods can be passed around and referenced like any other value in the language
- for example, say I that I want to customize the behavior of the `defaultdict` class
- this data structure allows you to supply a function that will be called with no arguments each time a missing key is accessed
- the function must return the default value that the missing key should have in the dictionary
- here I define a hook that logs each time a key is missing and returns 0 for the default value

```
[17]: def log_missing():
print('key added')
return 0
```

- given an initial dictionary and a set of desired increments, I can cause the `log_missing` function to run and print twice (for 'red' and 'orange')

```
[19]: from collections import defaultdict

current = {'green': 12, 'blue': 3}
increments = [
    ('red', 5),
    ('blue', 17),
    ('orange', 9),
]
```

```

result = defaultdict(log_missing, current)
print('Before:', dict(result))

for key, amount in increments:
    result[key] += amount
print('After: ', dict(result))

```

Before: {'green': 12, 'blue': 3}

key added

key added

After: {'green': 12, 'blue': 20, 'red': 5, 'orange': 9}

- supplying functions like `log_missing` makes the APIs easy to build and test because it separates side effects from deterministic behavior
- for example, say I want the default value hook passed to `defaultdict` to count the total number of keys that were missing
- one way to achieve this is by using a stateful closure (function inside of function)
- we define a helper function that uses such a closure as the default value hook

```

[20]: def increment_with_report(current, increments):
    added_count = 0

    def missing():
        nonlocal added_count # stateful closure
        added_count += 1
        return 0

    result = defaultdict(missing, current)
    for key, amount in increments:
        result[key] += amount

    return result, added_count

```

- running the function produces the expected result (2), even though the `defaultdict` has no idea that the `missing` hook maintains state
- another benefit of accepting simple functions for interfaces is that it's easy to add functionality later by hiding state in a closure

```

[21]: result, count = increment_with_report(current, increments)
assert count == 2

```

- the problem with defining a closure for stateful hooks is that it's harder to read than a stateless function
- another approach is to define a small class that encapsulates the state you want to track

```

[22]: class CountMissing:
    def __init__(self):
        self.added = 0

```

```
def missing(self):
    self.added += 1
    return 0
```

- in other languages, you might expect that now `defaultdict` would have to be modified to accommodate the interface of `CountMissing`
- but in Python, thanks to first-class functions, you can reference the `CountMissing.missing` directly on an object and pass it to `defaultdict` as the default value hook

```
[23]: counter = CountMissing()
result = defaultdict(counter.missing, current) # Method ref
for key, amount in increments:
    result[key] += amount
assert counter.added == 2
```

- using a helper class like this to provide the behavior of a stateful closure is cleaner than using the `increment_with_report` function
- in isolation however, it's still not clear what the purpose of the `CountMissing` class is
- who constructs the `CountMissing` object, who calls the `missing` method
- unless you see its usage with `defaultdict`, the class is a mystery
- to clarify the situation, Python allows classes to define the `__call__` special method
- `__call__` allows an object to be called just like a function
- it also causes the callable built-in function to return `True` for such instances, just like a normal function or method
- all objects that can be executed in this manner are referred to as `callables`

```
[24]: class BetterCountMissing:
    def __init__(self):
        self.added = 0

    def __call__(self):
        self.added += 1
        return 0

counter = BetterCountMissing()
assert counter() == 0
assert callable(counter)
```

- here we use a `BetterCountMissing` instance as the default value hook for a `defaultdict` to track the number of missing keys that were added

```
[25]: counter = BetterCountMissing()
result = defaultdict(counter, current) # Relies on __call__
for key, amount in increments:
    result[key] += amount
assert counter.added == 2
```

- this is much clearer than the `CountMissing.missing` example

- the `__call__` method indicates that a class's instances will be used somewhere a function argument would also be suitable (like API hooks)
- it directs new readers of the code to the entry point that's responsible for the class's primary behavior
- it provides a strong hint that the goal of the class is to act as a stateful closure
- best of all `defaultdict` still has no view into what's going on when you use `__call__`
- all that `defaultdict` requires is a function for the default value hook
- python provides different ways to satisfy a simple function interface and you can choose the one that works best for what you need to accomplish
- instead of defining and instantiating classes, you can often simply use functions for simple interfaces between components in python
- references to functions and methods in python are first class, meaning they can be used in expressions
- the `__call__` special method enables instances of a class to be called like plain python function
- when you need a function to maintain state, consider defining a class that provides the `__call__` method instead of defining a stateful closure

0.3 Item 39: Use @classmethod Polymorphism to Construct Objects Generically

- in python not only do objects support polymorphism, but classes do as well
- polymorphism enables multiple classes in a hierarchy to implement their own unique versions of a method
- this means that many classes can fulfill the same interface or abstract base class while providing different functionality
- for example, say that I'm writing a `MapReduce` implementation and I want a common class to represent the input data
- Here, I define such a class with a `read` method that must be defined by subclasses

```
[65]: class InputData:
    def __init__(self, data):
        self.data = data

    def read(self):
        raise NotImplementedError
```

- I also have a concrete subclass of `InputData` that reads data from a file on disk

```
[ ]: class PathInputData(InputData):
    def __init__(self, path):
        super().__init__()
        self.path = path

    def read(self):
        with open(self.path) as f:
            return f.read()
```

- you could have any number of `InputData` subclasses, like `PathInputData` and each of them could implement the standard interface for `read` to return the data to process

- other `InputData` subclasses could read the network, decompress data transparently, and so on
- I'd want a similar abstract interface for `MapReduce` worker that consumes the input data in a standard way

```
[15]: class Worker:
    def __init__(self, input_data):
        self.input_data = input_data
        self.result = None

    def map(self):
        raise NotImplementedError

    def reduce(self, other):
        raise NotImplementedError
```

- below I define a concrete subclass of `Worker` to implement the specific `MapReduce` function I want to apply- a simple newline counter

```
[16]: class LineCountWorker(Worker):
    def map(self):
        data = self.input_data.read()
        self.result = data.count('\n')

    def reduce(self, other):
        self.result += other.result
```

- one issue we have is that all of the classes are good but what connects?
- the classes are only useful once the objects are constructed
- what is responsible for building the objects and orchestrating the `MapReduce`?
- the simplest approach is to manually build and connect the objects with some helper functions
- below we list the contents of a directory and construct a `PathInputData` instance for each file it contains

```
[17]: import os

def generate_inputs(data_dir):
    for name in os.listdir(data_dir):
        yield PathInputData(os.path.join(data_dir, name))
```

- next I create the `LineCountWorker` instances by using the `InputData` instances returned by `generate_inputs`

```
[18]: def create_workers(input_list):
    workers = []
    for input_data in input_list:
        workers.append(LineCountWorker(input_data))
    return workers
```

- we execute these `Worker` instances by fanning out the `map` step to multiple threads
- then we call `reduce` repeatedly to combine the results into one final value

```
[56]: from threading import Thread

def execute(workers):
    threads = [Thread(target=w.map) for w in workers]
    for thread in threads: thread.start()
    for thread in threads: thread.join()

    first, *rest = workers
    for worker in rest:
        first.reduce(worker)
    return first.result
```

- finally we connect all the pieces together in a function to run each step

```
[57]: def mapreduce(data_dir):
    inputs = generate_inputs(data_dir)
    workers = create_workers(inputs)
    return execute(workers)
```

- running this function on a set of test input files works as expected

```
[58]: import os
import random

def write_test_files(tmpdir):
    os.makedirs(tmpdir)
    for i in range(100):
        with open(os.path.join(tmpdir, str(i)), 'w') as f:
            f.write('\n' * random.randint(0, 100))

# tmpdir = 'test_inputs' + str(random.randint(1, 101))
# write_test_files(tmpdir)

# result = mapreduce(tmpdir)
# print(f'There are {result} lines')
```

Problem with the code above: - the huge issue is that the `mapreduce` function is not generic at all - If we wanted to write another `InputData` or `Worker` subclass, I would also have to rewrite the `generate_input`, `create_workers` and `mapreduce` functions to match

- the problem boils down to needing a generic way to construct objects
- in other programming languages, you could use **constructor polymorphism**
 - you would require that each `InputData` subclass provides a special constructor that can be used generically by the helper methods that orchestrate the `MapReduce` (similar to the factory pattern)
- problem in python is that it only allows for a single constructor method `__init__`
- it's unreasonable to require every `InputData` subclass to have a compatible constructor

- the best way to solve this problem is with a `class method` polymorphism
- this is exactly like the instance method polymorphisms used for `InputData.read`, except that it's for whole classes instead of their constructed objects
- lets apply this idea to the `MapReduce` classes
- we can extend the `InputData` class with a generic `@classmethod` that's responsible for creating new `InputData` instance using a common interface

```
[59]: class GenericInputData:
    def read(self):
        raise NotImplementedError

    @classmethod
    def generate_inputs(cls, config):
        raise NotImplementedError
```

- we have `generate_inputs` take a dictionary with a set of configuration parameters that the `GenericInputData` concrete subclass needs to interpret
- below I use the `config` to find the directory to list for input files

```
[60]: class PathInputData(GenericInputData):
    def __init__(self, path):
        super().__init__()
        self.path = path

    def read(self):
        with open(self.path) as f:
            return f.read()

    @classmethod
    def generate_inputs(cls, config):
        data_dir = config['data_dir']
        for name in os.listdir(data_dir):
            yield cls(os.path.join(data_dir, name))
```

- similarly we can make the `create_workers` helper part of the `GenericWorker` class
- we use the `input_class` parameter, which must be a subclass of `GenericInputData`, to generate the necessary inputs
- I construct instances of the `GenericWorker` concrete subclass by using `cls()` as a generic constructor

```
[61]: class GenericWorker:
    def __init__(self, input_data):
        self.input_data = input_data
        self.result = None

    def map(self):
        raise NotImplementedError

    def reduce(self, other):
```

```

        raise NotImplementedError

    @classmethod
    def create_workers(cls, input_class, config):
        workers = []
        for input_data in input_class.generate_inputs(config):
            workers.append(cls(input_data))
        return workers

```

- note that the call to the `input_class.generate_inputs` above is the class polymorphism that were using to show
- you can see how `create_workers` calling `cls()` provides an alternative way to construct `GenericWorkers` object besides using the `__init__` method directly
- the effect on my concrete `GenericWorker` subclass is nothing more than changing its parent class

```

[62]: class LineCountWorker(GenericWorker):
        def map(self):
            data = self.input_data.read()
            self.result = data.count('\n')

        def reduce(self, other):
            self.result += other.result

```

finally we can rewrite the `mapreduce` function to be completely generic by calling `create_workers`

```

[63]: def mapreduce(worker_class, input_class, config):
        workers = worker_class.create_workers(input_class, config)
        return execute(workers)

```

- running the new worker on a set of test files produces the same result as the old implementation
- the difference is that `mapreduce` function requires more parameters so that it can operate generically

```

[64]: config = {'data_dir': tmpdir}
result = mapreduce(LineCountWorker, PathInputData, config)
print(f'There are {result} lines')

```

There are 4910 lines

- python only supports a single constructor per class: the `__init__` method
- use `@classmethod` to define alternative constructors for your classes
- use class method polymorphism to provide generic ways to build and connect many concrete subclasses

0.4 Item 40: Initialize Parent Classes with `super`

- the old, simple way to initialize a parent class from a child class is to directly call the parent class's `__init__` method with the child instance

```
[3]: class MyBaseClass:
    def __init__(self, value):
        self.value = value

class MyChildClass(MyBaseClass):
    def __init__(self):
        MyBaseClass.__init__(self, 5)
```

- the code works well for basic class hierarchies but brakes in many cases
- if a class is affected by multiple inheritance (something to avoid in general), calling the superclasses `__init__` methods directly can lead to unpredictable behavior
- one problem is that the `__init__` call order isn't specified across all subclasses
- below, we define two parent classes that operate on the instance's `value` field

```
[6]: class TimesTwo:
    def __init__(self):
        self.value *= 2

class PlusFive:
    def __init__(self):
        self.value += 5

# this class defines its parent classes in one ordering
class OneWay(MyBaseClass, TimesTwo, PlusFive):
    def __init__(self, value):
        MyBaseClass.__init__(self, value)
        TimesTwo.__init__(self)
        PlusFive.__init__(self)

foo = OneWay(5)
print('First ordering value is (5 * 2) + 5 = ', foo.value)
```

First ordering value is (5 * 2) + 5 = 15

- we can define another class that defines the same parent classes but in a different ordering (PlusFive second in the `__init__`)
- we still kept the `TimesTwo.__init__` and `PlusFive.__init__` in the same order
- the classe's behavior doesnt match the order of the parent classes
- the conflict here between the inheritance base classes and the `__inti__` calls is hard to spot, which makes this espically difficult for new readers of the code to understand

```
[7]: class AnotherWay(MyBaseClass, PlusFive, TimesTwo):
    def __init__(self, value):
        MyBaseClass.__init__(self, value)
        TimesTwo.__init__(self)
        PlusFive.__init__(self)

bar = AnotherWay(5)
```

```
print('Second ordering value is ', bar.value)
```

Second ordering value is 15

- Another problem occurs with diamond inheritance
- Diamond Inheritance happens when a subclass inherits from two separate classes that have the same superclass somewhere in the hierarchy
- Diamond inheritance causes the common superclass's `__init__` method to run multiple times, causing unexpected behavior

```
[18]: class TimesSeven(MyBaseClass):
        def __init__(self, value):
            MyBaseClass.__init__(self, value)
            self.value *= 7

        class PlusNine(MyBaseClass):
            def __init__(self, value):
                MyBaseClass.__init__(self, value)
                self.value += 9
```

- then we define a child class that inherits from both of these classes, making `MyBaseClass` the top of the diamond

```
[19]: class ThisWay(TimesSeven, PlusNine):
        def __init__(self, value):
            TimesSeven.__init__(self, value)
            PlusNine.__init__(self, value)

foo = ThisWay(5)
print('Should be (5 * 7) + 9 = 44 but is', foo.value)
```

Should be (5 * 7) + 9 = 44 but is 14

- the call on the second parent class's constructor, `PlusNine.__init__`, causes `self.value` to be reset back to 5 when `MyBaseClass.__init__` gets called a second time
- that results in the calculation of `self.value` to be $5 + 9 = 14$, completely ignoring the effect of the `TimesSeven.__init__` constructor
- this behavior is surprising and can be very difficult to debug in more complex cases
- to solve these problems, Python has the `super` built-in function and standard method resolution order (MRO)
- `super` ensures that common superclasses in diamond hierarchies are run only once
- the MRO defines the ordering in which superclasses are initialized, following an algorithm called C3 linearization
- below we create a diamond-shaped class hierarchy again, but this time we use `super` to initialize the parent class

```
[20]: class TimesSevenCorrect(MyBaseClass):
        def __init__(self, value):
            super().__init__(value)
```

```

        self.value *= 7

class PlusNineCorrect(MyBaseClass):
    def __inti__(self, value):
        self.value += 9

```

- Now, the top part of the diamond, `MyBaseClass.__init__` is run only a single time
- the other parent classes are run in the order specified in the `class` statement

```

[24]: class GoodWay(TimesSevenCorrect, PlusNineCorrect):
        def __init__(self, value):
            super().__init__(value)

foo = GoodWay(5)
print('Should be 7 (5 _ 9) = 98 and is', foo.value)

```

Should be 7 (5 _ 9) = 98 and is 35

- this order may seem backwards at first meaning, `TimesSevenCorrect.__init__` should have run first
- the result should be $(5 * 7) + 9 == 44$ but it is not
- the ordering matched what the MRO defines for the class
- The MRO ordering is available on a class method called `mro`

```

[25]: mro_str = '\n'.join(repr(cls) for cls in GoodWay.mro())
print(mro_str)

```

```

<class '__main__.GoodWay'>
<class '__main__.TimesSevenCorrect'>
<class '__main__.PlusNineCorrect'>
<class '__main__.MyBaseClass'>
<class 'object'>

```

- when we call `GoodWay(5)` it in turn calls `TimesSevenCorrect.__init__`, which calls `PlusNineCorrect.__init__`, which calls `MyBaseClass.__init__`
- once this reaches the top of the diamond, all of the initialization methods actually do their work in the opposite order from how their `__init__` functions were called
- Besides making multiple inheritance robust, the call to `super().__init__` is also much more maintainable than calling `MyBaseClass.__init__` directly in the subclasses
- I could later rename `MyBaseClass` to something else or have `TimesSevenCorrect` and `PlusNineCorrect` inherit from another superclass without having to update their `__init__` methods to match
- the `super` function can also be called with two parameters, first the type of the class whose MRO parent view you're trying to access and then the instance on which to access that view
- using these optional parameters within the constructor looks like:

```

[26]: class ExplicitTrisect(MyBaseClass):
        def __init__(self, value):
            super(ExplicitTrisect, self).__init__(value)

```

```
self.value /= 3
```

- those parameters are not needed as python's compiler can do it for you

```
[27]: class AutomaticTrisect(MyBaseClass):
    def __init__(self, value):
        super(__class__, self).__init__(value)
        self.value /= 3

class ImplicitTrisect(MyBaseClass):
    def __init__(self, value):
        super().__init__(value)
        self.value /= 3
```

```
[28]: assert ExplicitTrisect(9).value == 3
assert AutomaticTrisect(9).value == 3
assert ImplicitTrisect(9).value == 3
```

- the only time you should provide parameters to `super` is in situations where you need to access the specific functionality of a superclass's implementation from a child class (to wrap or reuse functionality)
- python's standard method resolution order (MRO) solves the problems of superclass initialization order and diamond inheritance
- use the `super` built-in function with zero arguments to initialize parent classes

0.5 Item 41: Consider Composing Functionality with Mix-in Classes

- python is an object-oriented language with built-in facilities for marking multiple inheritance tractable (`super`)
- however it's better to avoid multiple inheritance altogether
- if you want the convenience and encapsulation that comes with multiple inheritance, but want to avoid the potential headaches, consider writing a `mix-in` instead
- a `mix-in` is a class that defines only a small set of additional methods for its child classes to provide
- `mix-in` classes don't define their own instance attributes nor require their `__init__` constructor to be called
- writing `mix-ins` is easy because Python makes it trivial to inspect the current state of any object, regardless of its type
- Dynamic inspection means you can write generic functionality just once, in a `mix-in` and it can then be applied to many other classes
- `mix-ins` can be composed and layered to minimize repetitive code and maximize reuse
- let's say I want the ability to convert a Python object from its `in-memory` representation to a dictionary that's ready for serialization
- why not write this functionality generically so that I can use it with all my classes
- below we define an example `mixin` that accomplishes this with a new public method that's added to any class that inherits from it

- the implementation details are straightforward and rely on dynamic attribute access using `hasattr` dynamic type inspection with `isinstance` and accessing the instance dictionary `__dict__`

```
[40]: class ToDictMixin:
    def to_dict(self):
        return self._traverse_dict(self.__dict__)

    def _traverse_dict(self, instance_dict):
        output = {}
        for key, value in instance_dict.items():
            output[key] = self._traverse(key, value)
        return output

    def _traverse(self, key, value):
        if isinstance(value, ToDictMixin):
            return value.to_dict()
        elif isinstance(value, dict):
            return self._traverse_dict(value)
        elif isinstance(value, list):
            return [self._traverse(key, i) for i in value]
        elif hasattr(value, '__dict__'):
            return self._traverse_dict(value.__dict__)
        else:
            return value
```

- below we define an example class that uses the mix-in to make a dictionary representation of a binary tree

```
[41]: class BinaryTree(ToDictMixin):
    def __init__(self, value, left=None, right=None):
        self.value = value
        self.left = left
        self.right = right
```

- translating a large number of related Python objects into a dictionary becomes easy

```
[42]: tree = BinaryTree(10,
    left=BinaryTree(7, right=BinaryTree(9)),
    right=BinaryTree(13, left=BinaryTree(11)))

print(tree.to_dict())
```

```
{'value': 10, 'left': {'value': 7, 'left': None, 'right': {'value': 9, 'left':
None, 'right': None}}, 'right': {'value': 13, 'left': {'value': 11, 'left':
None, 'right': None}, 'right': None}}
```

- the best part about mix-ins is that you can make their generic functionality pluggable so behavior can be overridden when required
- for example, below I define a subclass of `BinaryTree` that hold a reference to its parent

- this circular reference would cause the default implementation of `ToDictMixin.to_dict` to loop forever
- the solution is to override the `BinaryTreeWithParent._traverse` method to only process values that matter, preventing cycles encountered by the mix-in
- the `_traverse` override inserts the parents numerical value and otherwise defers to the mix-ins default implementation by using the `super` built-functionality

```
[47]: class BinaryTreeWithParent(BinaryTree):
    def __init__(self, value, left=None,
                right=None, parent=None):
        super().__init__(value, left=left, right=right)
        self.parent = parent

    def _traverse(self, key, value):
        if (isinstance(value, BinaryTreeWithParent) and
            key == 'parent'):
            return value.value # Prevent cycles
        else:
            return super()._traverse(key, value)
```

- calling `BinaryTreeWithParent.to_dict` works without issue because the circular referencing properties aren't followed

```
[48]: root = BinaryTreeWithParent(10)
root.left = BinaryTreeWithParent(7, parent=root)
root.left.right = BinaryTreeWithParent(9, parent=root.left)
print(root.to_dict())
```

```
{'value': 10, 'left': {'value': 7, 'left': None, 'right': {'value': 9, 'left':
None, 'right': None, 'parent': 7}, 'parent': 10}, 'right': None, 'parent': None}
```

- by defining the `BinaryTreeWithParent._traverse`, I've also enabled any class that has an attribute type `BinaryTreeWithParent` to automatically work with the `ToDictMixin`

```
[49]: class NamedSubTree(ToDictMixin):
    def __init__(self, name, tree_with_parent):
        self.name = name
        self.tree_with_parent = tree_with_parent

my_tree = NamedSubTree('foobar', root.left.right)
print(my_tree.to_dict()) # No infinite loop
```

```
{'name': 'foobar', 'tree_with_parent': {'value': 9, 'left': None, 'right': None,
'parent': 7}}
```

- mix-ins can also be composed together
- let's say I want a mixin that provides generic JSON serialization for any class
- I can do this by assuming that a class provides a `to_dict` method (which may or may not be provided by the `ToDictMixin` class)

```
[56]: import json

class JsonMixin:
    @classmethod
    def from_json(cls, data):
        kwargs = json.loads(data)
        return cls(**kwargs)

    def to_json(self):
        return json.dumps(self.to_dict())
```

- note how the `JsonMixin` class defines both instance methods and class methods
- mix-ins let you add either kind of behavior to subclasses
- in this example, the only requirements of a `JsonMixin` subclass are providing a `to_dict` method and taking keyword arguments for the `__init__` method
- the mix-in makes it simple to create hierarchies of utility classes that can be serialized to and from JSON with little boilerplate
- for example, here I have a hierarchy of data classes representing parts of a datacenter topology

```
[57]: class DatacenterRack(ToDictMixin, JsonMixin):
    def __init__(self, switch=None, machines=None):
        self.switch = Switch(**switch)
        self.machines = [
            Machine(**kwargs) for kwargs in machines]

class Switch(ToDictMixin, JsonMixin):
    def __init__(self, ports=None, speed=None):
        self.ports = ports
        self.speed = speed

class Machine(ToDictMixin, JsonMixin):
    def __init__(self, cores=None, ram=None, disk=None):
        self.cores = cores
        self.ram = ram
        self.disk = disk
```

- serializing these classes to and from JSON is simple
- here we can verify that data is able to be sent round-trip through serializing and deserializing

```
[58]: serialized = """{
    "switch": {"ports": 5, "speed": 1e9},
    "machines": [
        {"cores": 8, "ram": 32e9, "disk": 5e12},
        {"cores": 4, "ram": 16e9, "disk": 1e12},
        {"cores": 2, "ram": 4e9, "disk": 500e9}
    ]
}"""
deserialized = DatacenterRack.from_json(serialized)
```

```
roundtrip = serialized.to_json()
assert json.loads(serialized) == json.loads(roundtrip)
```

- when you use `mix-ins` like this, its fine if the class you apply `JsonMixin` to already inherit from `JsonMixin` higher up the class hierarchy
- the resulting class will behave the same way, thanks to the behavior of `super`

0.5.1 Things to Remember

- avoid using multiple inheritance with instance attributes and `__init__` if mix-in classes can achieve the same outcome
- use pluggable behaviors at the instance level to provide per-class customization with `mix-in` classes may require
- `mix-ins` can include instance methods or class methods, depending on your needs
- compose `mix-ins` to create complex functionality from simple behaviors

0.6 Item 42: Prefer Public Attributes Over Private Ones

- there are only two types of visibility for a class's attribute: `public` or `private`

```
[1]: class MyObject:
      def __init__(self):
          self.public_field = 5
          self.__private_field = 10

      def get_private_field(self):
          return self.__private_field
```

- public attributes can be accessed by anyone using the dot operator on the object

```
[2]: foo = MyObject()
      assert foo.public_field == 5
```

- private fields are specified by prefixing an attribute's name with a double underscore
- they can be accessed directly by methods of the containing class

```
[3]: assert foo.get_private_field() == 10
```

- however directly accessing private fields from outside the class raises an exception:

```
[4]: foo.__private_field
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-4-a888a87e4048> in <module>
----> 1 foo.__private_field

AttributeError: 'MyObject' object has no attribute '__private_field'
```

- class methods also have access to private attributes because they are declared within the surrounding class block

```
[6]: class MyOtherObject:
    def __init__(self):
        self.__private_field = 71

    @classmethod
    def get_private_field_of_instance(cls, instance):
        return instance.__private_field

bar = MyOtherObject()
assert MyOtherObject.get_private_field_of_instance(bar) == 71
```

- as you'd expect with private fields, a subclass can't access its parent class's private fields

```
[9]: class MyParentObject:
    def __init__(self):
        self.__private_field = 71

class MyChildObject(MyParentObject):
    def get_private_field(self):
        return self.__private_field

baz = MyChildObject()
baz.get_private_field()
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-9-38ca44b74faf> in <module>
      8
      9 baz = MyChildObject()
----> 10 baz.get_private_field()

<ipython-input-9-38ca44b74faf> in get_private_field(self)
      5 class MyChildObject(MyParentObject):
      6     def get_private_field(self):
----> 7         return self.__private_field
      8
      9 baz = MyChildObject()

AttributeError: 'MyChildObject' object has no attribute
↳ '_MyChildObject__private_field'
```

- accessing the parents private attribute from the child class fails simply because the transformed attribute name doesn't exist
- knowing the schema, you can easily access the private attributes of any class- from a subclass or externally- without asking for permission

```
[11]: assert baz._MyParentObject__private_field == 71
```

- if you look in the object's attribute dictionary, you can see the private attributes are actually stored with the names as they appear after the transformation

```
[12]: print(baz.__dict__)
```

```
{'_MyParentObject__private_field': 71}
```

- python does not enforce not using private attributes because it believes that “we are all consenting adults here”
- PEP8 specifies that fields prefixed by a single underscore, `_protected_field` are protected by convention
- new programmers use `__` but should not because when someone wants to use the private attribute they will find a way
- document each protected field and explain which fields are internal APIs available to subclasses and which should be left out entirely

```
[13]: class MyStringClass:
    def __init__(self, value):
        # This stores the user-supplied value for the object
        # It should be coercible to a string. Once assigned in
        # the object it should be treated as immutable
        self._value = value
```

- the only time to seriously consider using private attributes is when you're worried about naming conflicts with subclasses
- this problem occurs when a child class unwittingly defines an attribute that was already defined by its parent class

```
[16]: class ApiClass:
    def __init__(self):
        self._value = 5

    def get(self):
        return self._value

class Child(ApiClass):
    def __init__(self):
        super().__init__()
        self._value = 'hello'

a = Child()
print(f'{a.get()} and {a._value} should be different')
```

hello and hello should be different

- this is primarily a concern with classes that are part of a public API
- the subclasses are out of your control, so you can't refactor to fix the problem
- such a conflict is especially possible with attribute names that are very common (like `value`)

- to reduce the risk of this issue occurring, you can use a private attribute in the parent class to ensure that there are no attribute names that overlap with child classes

```
[17]: class ApiClass:
    def __init__(self):
        self.__value = 5      # Double underscore

    def get(self):
        return self.__value  # Double underscore

class Child(ApiClass):
    def __init__(self):
        super().__init__()
        self._value = 'hello'

a = Child()
print(f'{a.get()} and {a._value} are different')
```

5 and hello are different

0.6.1 Things to Remember

- private attributes aren't rigorously enforced by the Python compiler
- Plan from the beginning to allow subclasses to do more with your internal APIs and attributes instead of choosing to lock them out
- use documentation of protected fields to guide subclasses instead of trying to force access control with private attributes
- only consider using private attributes to avoid naming conflicts with subclasses that are out of your control

0.7 Item 43: Inherit from `collections.abc` for Custom Container Types

- much of programming is defining classes that contain data and describing how such objects relate to each other
- every python class is a container of some kind, encapsulating attributes and functionality together
- python also provides built-in container types for managing data: `lists`, `tuples`, `sets` and `dictionaries`
- when you're designing classes for simple use cases like sequences, it's natural to want to subclass Python's built-in `list` type directly
- for example, say I want to create my own custom `list` type that has additional methods for counting the frequency of its members

```
[18]: class FrequencyList(list):
    def __init__(self, members):
        super().__init__(members)

    def frequency(self):
```

```

counts = {}
for item in self:
    counts[item] = counts.get(item, 0) + 1
return counts

```

- by subclassing `list`, I get all of `list`'s standard functionality and preserve the semantics familiar to Python programmers
- we can define additional methods to provide any custom behaviors that I need

```

[22]: foo = FrequencyList(['a', 'b', 'c', 'b', 'a', 'd'])
print('Length is', len(foo))

foo.pop()
print('After pop:', repr(foo))
print('Frequency:', foo.frequency())

```

Length is 6

After pop: ['a', 'b', 'c', 'b', 'a']

Frequency: {'a': 2, 'b': 2, 'c': 1}

- now imagine that I want to provide an object that feels like a `list` and allows indexing but is not a `list` subclass
- for example, say that I want to provide sequence semantics (like `list` or `tuple` for a binary tree class

```

[24]: class BinaryNode:
    def __init__(self, value, left=None, right=None):
        self.value = value
        self.left = left
        self.right = right

```

- how do you make this class act like a sequence type? Python implements its container behaviors with instance methods that have special names

```

[25]: bar = [1, 2, 3]
bar[0]

```

[25]: 1

- it will be interpreted as:
 - `bar.__getitem__(0)`
- to make the `BinaryNode` class act like a sequence, you can provide a custom implementation of `__getitem__` that traverses the object tree depth first

```

[32]: class IndexableNode(BinaryNode):
    def _traverse(self):
        if self.left is not None:
            yield from self.left._traverse()
        yield self

```



```

        if self.right is not None:
            yield from self.right._traverse()

    def __getitem__(self, index):
        for i, item in enumerate(self._traverse()):
            if i == index:
                return item.value
        raise IndexError(f'Index {index} is out of range')

tree = IndexableNode(
    10,
    left=IndexableNode(
        5,
        left=IndexableNode(2),
        right=IndexableNode(
            6,
            right=IndexableNode(7))),
    right=IndexableNode(
        15,
        left=IndexableNode(11)))

```

- but you can also access it like a list in addition to being able to traverse the tree with the left and right attributes

```

[33]: print('LRR is', tree.left.right.right.value)
      print('Index 0 is', tree[0])
      print('Index 1 is', tree[1])
      print('11 in the tree?', 11 in tree)
      print('17 in the tree?', 17 in tree)
      print('Tree is', list(tree))

```

```

LRR is 7
Index 0 is 2
Index 1 is 5
11 in the tree? True
17 in the tree? False
Tree is [2, 5, 6, 7, 10, 11, 15]

```

- the problem is implementing `__getitem__` isn't enough to provide all of the sequence semantics you'd expect from a list instance
- the `len` built-in function requires another special method named `__len__`, that must have an implementation for a custom sequence type

```

[34]: class SequenceNode(IndexableNode):
      def __len__(self):
          for count, _ in enumerate(self._traverse(), 1):
              pass
          return count

```

```
[35]: tree = SequenceNode(
    10,
    left=SequenceNode(
        5,
        left=SequenceNode(2),
        right=SequenceNode(
            6,
            right=SequenceNode(7))),
    right=SequenceNode(
        15,
        left=SequenceNode(11))
)

print('Tree length is', len(tree))
```

Tree length is 7

- unfortunately, this still isn't enough for the class to fully be a valid sequence
- also missing are the `count` and `index` methods that a Python programmer would expect to see on a sequence like `list` or `tuple`
- it turns out that defining your own container type is much harder than it seems
- to avoid this difficulty throughout the Python universe, the built-in `collections.abc` module defines a set of abstract base classes that provide all the typical methods for each container type
- when you subclass from these abstract base classes and forget to implement required methods, the module tells you something is wrong

```
[36]: from collections.abc import Sequence

class BadType(Sequence):
    pass

foo = BadType()
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-36-997f61651b50> in <module>
      4     pass
      5
----> 6 foo = BadType()

TypeError: Can't instantiate abstract class BadType with abstract methods
↳ __getitem__, __len__
```

- when you do implement all the methods required by an abstract base class from `collections.abc`, as done with the `SequenceNode`, it provides all of the additional methods like `index`, and `count` for free

```
[38]: class BetterNode(SequenceNode, Sequence):
        pass

tree = BetterNode(
    10,
    left=BetterNode(
        5,
        left=BetterNode(2),
        right=BetterNode(
            6,
            right=BetterNode(7))),
    right=BetterNode(
        15,
        left=BetterNode(11))
)

print('Index of 7 is', tree.index(7))
print('Count of 10 is', tree.count(10))
```

```
Index of 7 is 3
Count of 10 is 1
```

- the benefit of using these abstract base classes is even greater for more complex container types such as `Set` and `MutableMapping`, which have a large number of special methods that need to be implemented to match Python conventions
- beyond the `collections.abc` module, Python uses a variety of special methods for object comparison and sorting which may be provided by container classes and non-container classes
- inherit directly from Python's container types (like `list` or `dict`) for simple use cases
- beware of the large number of methods required to implement custom container types correctly
- have your custom container types inherit from the interfaces defined in `collections.abc` to ensure that your classes match required interface and behavior