

# Chapter 13 - Concurrency

April 4, 2021

## 0.1 Overview

- concurrency is the art of making a computer do multiple things at once
- concurrency concepts are fairly straight forward but the bugs that can occur are notoriously difficult to track down

## 0.2 Threads

- most often concurrency is created so that work can continue happening while the program is waiting for I/O to happen
- we can rely on python and the operating system to take care of the tricky switching part, while we create objects that appear to be running independently, by simultaneously
- these objects are called **threads**

```
[2]: from threading import Thread

class InputReader(Thread):
    def run(self):
        self.line_of_text = input()

print("Enter some text and press enter: ")
thread = InputReader()
thread.start()

count = result = 1
while thread.is_alive():
    result = count * count
    count += 1

print(f'calculated squares up to {count} * {count} = {result}')
print(f"while you typed {thread.line_of_text}")
```

Enter some text and press enter:

s

calculated squares up to 1991670 \* 1991670 = 3966745405561

while you typed s

- example above runs two threads
- every program has a single thread, here we introduced another thread called **InputReader**
- to construct a thread, we must extend the **Thread** class and implement the **run** method

- any code executed by the `run` method happens in a separate thread
- the new thread does not start until we call `.start()` method on the object
- if we want to take out the concurrent call to see how it compares, we can call `thread.run()` in the place that we originally called `thread.start()`
- interestingly, data we construct in one thread is accessible from other running threads
- remember, just because a method is on a `Thread` instance does not mean it is magically executed inside that thread
- there is a `thread.join()` method that says to *wait for the thread to complete before doing anything*

```
[ ]: # code below ensures threads won't close until all of them finish
for thread in threads:
    thread.join()
```

## 0.3 The Many Problems With Threads

### 0.3.1 Shared Memory

- the threads have access to all the program's memory and thus all its variables
- this can too easily cause inconsistencies in the program state
- the solution to this problem in threaded programs is to **synchronize** access to any code that reads or writes a shared variable
- synchronization works but you have to remember to try and use it and it introduces bugs that are hard to track down

### 0.3.2 Global Interpreter Lock

- in order to efficiently manage memory, garbage collection and calls to machine code, Python has a utility called the **global interpreter lock (GIL)**
- this means threads are useless in python for the one thing they excel in other languages: **parallel processing**
- GIL prevents any two threads from doing work at the exact same time
- work means using the CPU, they can however use API calls or read data from the disk

### 0.3.3 Thread Overhead

- there is also a cost of maintaining each thread
- each thread takes up a certain amount of memory to record the state of that thread
- switching between threads also uses some CPU time
- this can be solved by structuring our workload so that the threads can be reused to perform multiple jobs
- python provides a `ThreadPool` feature to handle this
- `ThreadPool` behaves similar to `ProcessPool`

## 0.4 Multiprocessing

- Multiprocessing library was designed to mimic the `thread` API but it has evolved to provide more robust features

- multiprocessing is not useful when the processes spend a majority of their time waiting on I/O but it is the way to go for parallel computation
- multiprocessing module spins up new operating system processes to do the work
- this means there is an entire python interpreter running each process

```
[4]: from multiprocessing import Process, cpu_count
import time
import os

class MuchCPU(Process):
    def run(self):
        print(os.getpid())
        for i in range(200000000):
            pass

if __name__ == "__main__":
    procs = [MuchCPU() for f in range(cpu_count())]
    t = time.time()
    for p in procs:
        p.start()
    for p in procs:
        p.join()
    print(f"work took {(time.time()) - t} seconds")
```

work took 0.2248551845550537 seconds

- each Process has a pid number which we can get using the `os.getpid()`
- also notice the `__name__` guardrail we put up so we dont accidentally import MuchCPU
- the difference between `Thead` and `Process` is that process is three times faster then thread

## 0.5 Multiprocessing Pools

for the following reasons, you should not have more processes than there are processors on the computer

- only `cpu_count()` process can run simultaneously
- each process consumes resources with a full copy of the Python interpreter
- communication between process is expensive
- crating process takes a non-zero amount of time
- you want to create at most `cpu_count()` processes when the program starts
- doing this on your own can be increadably difficult, but developers have done this for you in the form of `multiprocessing pools`
- `Pools` abstract away the overhead of figuring out what code is executing in the main process and which code is running in the subprocess
- the pool abstraction restricts number of places in which code in different processes interacts, making it much easier to keep track of
- unlike threads, multiprocessing cannot directly access variables set up by other threads
- pools hide the process of passing data between processes

- using a pool looks like a function call; you pass data into a function, it is executed in another process or processes, and when the work is done, a value is returned
- behind the scenes, there is a lot of work being done such as an object getting pickled and then passed, etc
- pickling takes a lot of time, and thus you should only pass a minimum amount of data between pools

```
[ ]: import random
from multiprocessing.pool import Pool

def prime_factor(value):
    factor = []
    for divisor in range(2, value - 1):
        quotient, remainder = divmod(value, divisor)
        if not remainder:
            factors.extend(prime_factor(divisor))
            factor.extend(prime_factor(quotient))
            break
    else:
        factors = [value]
    return factor

if __name__ == "__main__":
    pool = Pool()

    to_factor = [random.randint(10000, 5000000) for i in range(20)]
    results = pool.map(prime_factor, to_factor)
    for value, factors in zip(to_factor, results):
        print(f'the factors of {value} are {factors}')
```

- we first construct a multiprocessing pool instance
- by default this pool creates a separate process for each of the CPU cores in the machine
- the `map` method accepts a function and an iterable
- the pool pickles each of the values in the iterable and passes it into an available process, which executes the function on it
- when the process is finished doing its work, it pickles the resulting list of factors and passes it back to the pool
- then the pool has more work available, it takes on the next job
- there are funky `await/sync` stuff you can do with the pool
- you can also close a pool or terminate it

## 0.6 Queues

- if we need more control over communication between processes we can use a queue
- Queue data structures are useful for sending messages from one process into another processor
- any pickled object can be sent into a Queue
- this idea of a queue can actually become a distributed system

## 0.7 Problem with Multiprocessing

- primary drawback is that sharing data between processes is costly
- excessive pickling quickly dominates processing time
- multiprocessing works best when relatively small objects are passed between processes and a tremendous amount of work need to be done on each one
- the other major problem, is that it can be hard to tell which process a variable or method is being accessed

## 0.8 Futures

- there are also asynchronous ways of implementing concurrency
- **Futures** wrap either multiprocessing or threading depending on which concurrency we need
- whether you need I/O (threading) or CPU (multiprocessing)
- **Futures** are useful for *call and answer* type interactions
- meaning, processing can happen in another thread and at some point in the future, you can ask it for the results
- futures are just a wrapper around multiprocessing pools and thread pools

```
[6]: from concurrent.futures import ThreadPoolExecutor
from pathlib import Path
from os.path import sep as pathsep
from collections import deque

def find_files(path, query_string):
    subdirs = []
    for p in path.iterdir():
        full_path = str(p.absolute())
        if p.is_dir() and not p.is_symlink():
            subdirs.append(p)
        if query_string in full_path:
            print(full_path)
    return subdirs

query = '.py'
futures = deque()
basedir = Path(pathsep).absolute()

with ThreadPoolExecutor(max_workers=10) as executor:
    futures.append(
        executor.submit(find_files, basedir, query))
    while futures:
        future = futures.popleft()
        if future.exception():
            continue
        elif future.done():
            subdirs = future.result()
            for subdir in subdirs:
```

```

        futures.append(executor.submit(
            find_files, subdir, query))
    else:
        futures.append(future)

```

```

C:\$Recycle.Bin\S-1-5-21-2588907532-1068130059-599424490-1001\I07ZBC4.py
C:\Users\Vicktree\Desktop\test.py.un~
C:\Users\Vicktree\Desktop\test.py~.un~
C:\Users\Vicktree\Desktop\test.py
C:\Users\Vicktree\Desktop\test.py~
C:\Users\Vicktree\Downloads\password_ongoing.py

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-6-d584e51053b2> in <module>
    24         future = futures.popleft()
--> 25         if future.exception():
    26             continue

c:
↳ \users\vicktree\appdata\local\programs\python\python39\lib\concurrent\futures_base.
↳ py in exception(self, timeout)
    467
--> 468         self._condition.wait(timeout)
    469

c:\users\vicktree\appdata\local\programs\python\python39\lib\threading.py in
↳ wait(self, timeout)
    311         if timeout is None:
--> 312             waiter.acquire()
    313             gotit = True

```

KeyboardInterrupt:

During handling of the above exception, another exception occurred:

```

KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-6-d584e51053b2> in <module>
    31             find_files, subdir, query))
    32         else:
--> 33         futures.append(future)

c:
↳ \users\vicktree\appdata\local\programs\python\python39\lib\concurrent\futures_base.
↳ py in __exit__(self, exc_type, exc_val, exc_tb)
    626
    627     def __exit__(self, exc_type, exc_val, exc_tb):
--> 628         self.shutdown(wait=True)

```

```

629         return False
630
c:
↪ \users\vicktrees\appdata\local\programs\python\python39\lib\concurrent\futures\thread.
↪ py in shutdown(self, wait, cancel_futures)
    227         if wait:
    228             for t in self._threads:
--> 229                 t.join()
    230     shutdown.__doc__ = _base.Executor.shutdown.__doc__

c:\users\vicktrees\appdata\local\programs\python\python39\lib\threading.py in
↪ join(self, timeout)
    1031
    1032     if timeout is None:
-> 1033         self._wait_for_tstate_lock()
    1034     else:
    1035         # the behavior of a negative timeout isn't documented, but

c:\users\vicktrees\appdata\local\programs\python\python39\lib\threading.py in
↪ _wait_for_tstate_lock(self, block, timeout)
    1047     if lock is None: # already determined that the C code is done
    1048         assert self._is_stopped
-> 1049     elif lock.acquire(block, timeout):
    1050         lock.release()
    1051     self._stop()

KeyboardInterrupt:

```

- the core of the program above is the event loop
- we can construct a `ThreadPoolExecutor` as a context manager so that it is automatically cleaned-up and closes its threads when it is done
- it requires a `max_workers` argument to indicate the number of threads running at the time
- `ProcessPoolExecutor` normally is constrained to the number of CPU's on the machine, but with threads, it can be much higher, depending how many are waiting on I/O at the time
- once the executor had been constructed, we submit a job to it using the `submit` method
- the `submit()` method immediately returns a `Future` object, which promises to give us a result
- the future is placed inside the queue
- the loop then repeatedly removes the first future from the queue and inspects it
- if it is still running, it gets added back to the end of the queue
- if no errors occur, we can call `result()` to get the return value

## 0.9 AsyncIO

- combines the concepts of futures and event loops with the `coroutines`
- this was specifically designed for network I/O
- this library provides its own event loop
- the cost of this event loop is that when we run code in the `async` task on the event loop, the

- code must return immediately
- blocking neither on I/O nor on long-running calculation
- AsyncIO solves this by creating a set of `coroutines` that use `async` and `await` syntax to return control to the event loop immediately when the code will block
- the keyword replaces the `yield`, `yield from` and `send` syntax we used with `raw` `coroutines`

## 0.10 AsyncIO in Action

```
[8]: import asyncio
import random

async def random_sleep(counter):
    delay = random.random() * 5
    print(f"{counter} sleeps for {delay}")
    await asyncio.sleep(delay)
    print(f"{counter} awakens")

async def five_sleepers():
    print("creating five tasks")
    tasks = [asyncio.create_task(random_sleep(i)) for i in range(5)]
    print("Sleeping after starting five tasks")
    await asyncio.sleep(2)
    print("Waking and waiting for five tasks")
    await asyncio.gather(*tasks)

asyncio.get_event_loop().run_until_complete(five_sleepers())
print("Done five tasks")
```

```
-----
RuntimeError                                Traceback (most recent call last)
<ipython-input-8-efcc171606e9> in <module>
    16     await asyncio.gather(*tasks)
    17
--> 18 asyncio.get_event_loop().run_until_complete(five_sleepers())
    19 print("Done five tasks")

c:\users\vicktreet\appdata\local\programs\python\python39\lib\asyncio\base_event_l
↳py in run_until_complete(self, future)
    616         """
    617         self._check_closed()
--> 618         self._check_running()
    619
    620         new_task = not futures.isfuture(future)

c:\users\vicktreet\appdata\local\programs\python\python39\lib\asyncio\base_event_l
↳py in _check_running(self)
    576     def _check_running(self):
```



```

577         if self.is_running():
--> 578             raise RuntimeError('This event loop is already running')
579         if events._get_running_loop() is not None:
580             raise RuntimeError(

```

**RuntimeError:** This event loop is already running

- A task in this context, is an object that `asyncio` knows how to schedule on the event loop

**This includes:** - coroutines defined with the `async` and `await` syntax - coroutines decorated with the `@asyncio.coroutine` and using the `yield from` syntax (deprecated in favor of the `async` and `await` - `asyncio.Future` objects. These are almost identical to the `concurrent.futures` but for the use with `asyncio` - any awaitable object, that is, one with an `__await__` function

**examining the `five_sleepers` future:**

- the coroutine first constructs five instances of the `random_sleep` coroutine
- these are wrapped in a `asyncio.create_task` call, which adds the future to the loop's task queue so they can execute and start immediately, when control is returned to the loop
- the control is returned whenever we call `await`
- in this case, we call `await asyncio.sleep` to pause the execution of the coroutine for two seconds
- during the break, the event loop executes the tasks that it has queue up, namely the five `random_sleep` tasks
- when the sleep call in the `five_sleepers` task wakes up, it calls `asyncio.gather`. this function accepts tasks as `varargs` and awaits each of them before returning
- each of the `random_sleep` coroutines prints a starting message, then sends control back to the event loop for a specific amount of time using its own `await` call
- when the sleep has completed, the event loop passes control back to the relevant `random_sleep` tasks, which prints its awakening message before returning
- when the event queue is empty, the `run_until_complete` call is able to terminate and the program ends
- `async` keyword acts as documentation notifying the python interpreter that the coroutine contains the `await` calls

## 0.11 Reading an AsyncIO Future

- an AsyncIO coroutine executes each line in order until it encounters an `await` statement at which point, it returns control to the event loop
- the event loop then executes any other tasks that are ready to run, including the one that the original coroutine was waiting on
- whenever that child task completes, the event loop sends the result back into the coroutine so that it can pick up execution until it encounters another `await` statement or returns

## 0.12 AsyncIO for Networks

- AsyncIO was specifically designed for use with network sockets
- AsyncIO networking revolves around the intimately linked concepts of transporting and protocols

- a protocol is a class that has specific methods that are called when relevant events happen
- Since DNS runs on top of UDP (**User Datagram Protocol**), we build our protocol class as a subclass of `DatagramProtocol`
- the transport essentially represents a communication stream
- behind the scenes the transport has set up a task on the event loop that is listening for incoming UDP connections
- all we have to do, then is start the event loop running with the call to `loop.run_forever()` so that the task can process these packets
- when the packets arrive, they are processed on the protocol and everything just works fine
- there is a lot of boilerplate in setting up a protocol class and the underlying transport
- AsyncIO provides an abstraction on top of these two key concepts, called **streams**
- we will see an example of streams in the TCP server

### 0.13 Using Executors to Wrap Blocking Code

- AsyncIO provides its own version of the futures library to allow us to run code in a separate thread or process when there is not an appropriate non-blocking call to be made
- this allows us to combine threads and processes with the asynchronous model
- this is useful when an application has bursts of I/O bound and CPU-bound activity

### 0.14 Streams

- below is description of code I chose not to include due to its verbosity
- `create_server` hooks into AsyncIO's streams instead of using the underlying transport/protocol code
- it allows us to pass in a normal coroutine, which receives reader and writer parameters
- these both represent streams of bytes that can be read from and written, like files or sockets
- secondly, because this is a TCP server, instead of **UDP**, there is some socket cleanup required
- this cleanup is a blocking call, so we have to run `wait_closed` coroutine on the event loop
- Streams reading is a potentially blocking call so we have to call it with `await`
- writing doesn't block; it just puts the data in queue, which AsyncIO sends out in the background

### 0.15 AsyncIO clients

- because it can handle many thousands of simultaneous connections, AsyncIO is a very common for implementing servers
- clients can be much simpler than servers, as they don't have to be set up to wait for incoming connections
- like most networking libraries, you just open a connection, submit your requests and process any responses
- the main difference is that you need to use `await` any time you make a potentially blocking call

```
[9]: import asyncio
import random
import json
```

```

async def remote_sort():
    reader, writer = await asyncio.open_connection("127.0.0.1", 2015)
    print("Generating random list...")
    numbers = [random.randrange(10000) for r in range(10000)]
    data = json.dumps(numbers).encode()
    print("List Generated, Sending data")
    writer.write(len(data).to_bytes(8, "big"))
    writer.write(data)

    print("Waiting for data...")
    data = await reader.readexactly(len(data))
    print("Received data")
    sorted_values = json.loads(data.decode())
    print(sorted_values)
    print("\n")
    writer.close()

loop = asyncio.get_event_loop()
loop.run_until_complete(remote_sort())
loop.close()

```

```

-----
RuntimeError                                Traceback (most recent call last)
<ipython-input-9-ee1349fe2c32> in <module>
    20
    21 loop = asyncio.get_event_loop()
--> 22 loop.run_until_complete(remote_sort())
    23 loop.close()
    24

c:\users\vicktrees\appdata\local\programs\python\python39\lib\asyncio\base_event_loop.py in run_until_complete(self, future)
    616         """
    617         self._check_closed()
--> 618         self._check_running()
    619
    620         new_task = not futures.isfuture(future)

c:\users\vicktrees\appdata\local\programs\python\python39\lib\asyncio\base_event_loop.py in _check_running(self)
    576     def _check_running(self):
    577         if self.is_running():
--> 578             raise RuntimeError('This event loop is already running')
    579         if events._get_running_loop() is not None:
    580             raise RuntimeError(

```

```
RuntimeError: This event loop is already running
```