SQLAlchemy Core

March 31, 2021

0.1 Init

```
[2]: import sys
     !{sys.executable} -m pip install SQLAlchemy
     import sqlalchemy
    WARNING: You are using pip version 20.2.3; however, version 21.0.1 is available.
    You should consider upgrading via the
    'c:\users\vicktree\appdata\local\programs\python\python39\python.exe -m pip
    install --upgrade pip' command.
    WARNING: You are using pip version 20.2.3; however, version 21.0.1 is available.
    You should consider upgrading via the
    'c:\users\vicktree\appdata\local\programs\python\python39\python.exe -m pip
    install --upgrade pip' command.
    Collecting SQLAlchemy
      Downloading SQLAlchemy-1.4.4-cp39-cp39-win_amd64.whl (1.5 MB)
    Collecting greenlet!=0.4.17; python_version >= "3"
      Downloading greenlet-1.0.0-cp39-cp39-win_amd64.whl (95 kB)
    Installing collected packages: greenlet, SQLAlchemy
    Successfully installed SQLAlchemy-1.4.4 greenlet-1.0.0
    Collecting SQLAlchemy
      Downloading SQLAlchemy-1.4.4-cp39-cp39-win_amd64.whl (1.5 MB)
    Collecting greenlet!=0.4.17; python version >= "3"
      Downloading greenlet-1.0.0-cp39-cp39-win_amd64.whl (95 kB)
    Installing collected packages: greenlet, SQLAlchemy
    Successfully installed SQLAlchemy-1.4.4 greenlet-1.0.0
```

0.2 Schema and Types

- inorder to provide access to the underlying database, SQLALchemy needs a representation of the tables that should be present in the database
- this can be done in three ways
 - using user-defined Table objects
 - using declarative classes that represent your tables
 - interferring them from the database
- we will be using user-defined Table objects

• the 'Table object contain a list of typed columns and their attributes, which are associated with a common metadata container

0.2.1 Types

- there are four categories of types we can use inside of SQLAlchemy
- Generic
- Vendor specific
- User defined

Table 1-1. Generic type representations

SQLAIchemy	Python	SQL
BigInteger	int	BIGINT
Boolean	bool	BOOLEAN or SMALLINT
Date	datetime.date	DATE (SQLite: STRING)
DateTime	datetime.datetime	DATETIME (SQLite: STRING)
Enum	str	ENUM or VARCHAR
Float	float or Decimal	FLOAT or REAL
Integer	int	INTEGER
Interval	datetime.timedelta	INTERVAL or DATE from epoch
LargeBinary	byte	BLOB or BYTEA
Numeric	decimal.Decimal	NUMERIC or DECIMAL
Unicode	unicode	UNICODE or VARCHAR
Text	str	CLOB or TEXT
Time	datetime.time	DATETIME

- There are standard types which we can import from sqlchemy.dialects
- they are usually captial letters
- from sqlalchemy.dialects.postgresql iport JSON
- this allows us to use things like array_to_json
- we can also use VARCHAR when working with legacy codebase

0.2.2 Metadata

- Metadata is used to tie together the database structure so it can be quickly accessed inside SQLAlchemy
- metadata is a catalog of Table objects with optional information about the engine and connection
- these tables can be accessed via a dictionary, MetaData.tables

```
[2]: from sqlalchemy import MetaData
metadata = MetaData()
```

0.2.3 Tables

- Tables objects are initialized in SQLAlchemy Core in supplied MetaData object by calling the Table constructor with the table name and metadata
- any additional arguments are assumed to be column objects
- Column objects represent each field in the table
- Columns are constructed by calling Column with a name, type and then arguemnt that represent any additional SQL constructs and constraints

• note index(=True) is being used to speed up queires on that column

```
Column('updated_on', DateTime(), default=datetime.now, onupdate=datetime.

→now)
```

0.2.4 Keys and Constraints

- keys and constraints ensure that the data meets certain requirements prior to being stored in the database
- from sqlalchemy import PrimaryKeyConstraint, UniqueConstraint, CheckConstraints
- the most common key type is the primary key which is used to ensure a proper relationship between two pieces of related data in different tables
- if your primary_key=True on multiple columns, then the keys are treated as a tuple

•

- we can add multiple columns seprated by commas to create a composite key
- we can explicitly define a key with the following code:
 - PrimaryKeyConstraints('user_id', name='user_pk')
- another popular constraint is the unique constraint, which ensures that no two values are duplicated in a given field
 - UniqueConstraints('username', name='uix_username')
 - we are insuring that each customer had a unique username to log
- another one is the check constraint, which is used to ensure that the data supplied for a column matches a set of user-defined criteria
 - CheckConstraint('unit_cost >= 0.00', name='unit_cost_postive')

0.2.5 indexes

- used to accelerate lookups for field values
- we create index on cookie_name column because we know we will be searching by that often
- we can also define an index using an explicit construction type
- multiple columns can be designated by separating them by a comma
- you can also add a keyword argument of unique=True require the index to be unique

```
[5]: from sqlalchemy import Index

Index('ix_cookies_cookie_name', 'cookie _name')
```

- [5]: Index('ix_cookies_cookie_name', 'cookie _name')
 - we can also create functional indexes that vary a bit by the backend database being used
 - this lets you create an index for situations where you often need to query based on some unusual context
 - for example if we want to select by cookie SKU and name as a joined item, such as SKU001 Chocolate Chip
 - Index('ix test, mytable.c.cookie sku, mytable.c.cookie name)

0.2.6 Relationships and ForeignKeyConstraints

• one way to implement a relationshp is to use ForeignKeyConstraint

• in the ForeignKey you want to use string and not actual reference because using reference could actually cause it to fail

0.2.7 Persisting The Tables

- persisting the schema to the database is simply a matter of calling the create_all() method on our metadata instance with the engine where it should create those tables:
 - metadata.create_all(engine)
- better to use database migration tools like Alembic to handle any changes to existing tables or additional schema than to try to handcode changes directly in your application code

```
users = Table('users', metadata,
   Column('user_id', Integer(), primary_key=True),
   Column('customer_number', Integer(), autoincrement=True),
   Column('username', String(15), nullable=False, unique=True),
   Column('email_address', String(255), nullable=False),
   Column('phone', String(20), nullable=False),
   Column('password', String(25), nullable=False),
   Column('created_on', DateTime(), default=datetime.now),
   Column('updated on', DateTime(), default=datetime.now, onupdate=datetime.
→now)
orders = Table('orders', metadata,
   Column('order_id', Integer(), primary_key=True),
   Column('user_id', ForeignKey('users.user_id'))
line_items = Table('line_items', metadata,
   Column('line items id', Integer(), primary key=True),
   Column('order_id', ForeignKey('orders.order_id')),
   Column('cookie id', ForeignKey('cookies.cookie id')),
   Column('quantity', Integer()),
   Column('extended_cost', Numeric(12, 2))
engine = create_engine('sqlite:///:memory:')
metadata.create_all(engine)
```

0.3 Working With Data via SQLAlchemy Core

0.3.1 Inserting Data

```
print('')

connection = engine.connect()
result = connection.execute(ins)

print('ID of the recored we just created')
result.inserted_primary_key
```

The SQL equlivant of our insert statement
INSERT INTO cookies (cookie_name, cookie_recipe_url, cookie_sku, quantity,
unit_cost) VALUES (:cookie_name, :cookie_recipe_url, :cookie_sku, :quantity,
:unit_cost)

the paramaters entered into our SQL statement

ID of the recored we just created

[8]: [1]

- alternatively you could also insert by importing insert
- more adviced to just use the Table method insert

- - execute method of the connection object can take more than just statements
 - it is also possible to provide the values as keyword arguemnts to the execute method after our statement

```
[10]: ins = cookies.insert()
    result = connection.execute(
         ins,
            cookie_name='dark chocolate chip',
            cookie_recipe_url='https://some.aweso.me/cookie/recipe_dark.html',
            cookie_sku="CC02",
            quantity="1",
            unit_cost="0.75"
```

```
)
result.inserted_primary_key
```

[10]: [2]

• we can insert multiple objects into the table by using an array

```
[11]: inventory_list = [
       {
          'cookie_name': 'peanut butter',
          'cookie_recipe_url': 'http://some.aweso.me/cookie/peanut.html',
          'cookie sku': 'PB01',
          'quantity': '24',
          'unit_cost': '0.25'
       },
       {
          'cookie_name': 'oatmeal raisin',
          'cookie_recipe_url': 'http://some.okay.me/cookie/raisin.html',
          'cookie_sku': 'EWW01',
          'quantity': '100',
          'unit_cost': '1.00'
      }
      ]
      result = connection.execute(ins, inventory_list)
```

0.3.2 Querying Data

```
[12]: from sqlalchemy.sql import select

s = select([cookies])
rp = connection.execute(s)
results = rp.fetchall()

for cookie in results:
    print(f'{cookie} \n')
```

- (1, 'chocolate chip', 'http://some.aweso.me/cookie/recipe.html', 'CCO1', 12, Decimal('0.50'))
- (2, 'dark chocolate chip', 'https://some.aweso.me/cookie/recipe_dark.html', 'CCO2', 1, Decimal('0.75'))
- (3, 'peanut butter', 'http://some.aweso.me/cookie/peanut.html', 'PB01', 24, Decimal('0.25'))
- (4, 'oatmeal raisin', 'http://some.okay.me/cookie/raisin.html', 'EWW01', 100, Decimal('1.00'))

c:\users\vicktree\appdata\local\programs\python\python38-32\lib\sitepackages\sqlalchemy\sql\sqltypes.py:661: SAWarning: Dialect sqlite+pysqlite does
not support Decimal objects natively, and SQLAlchemy must convert from
floating point - rounding errors and other issues may occur. Please consider
storing Decimal numbers as strings or integers on this platform for lossless
storage.

util.warn(

```
[13]: from sqlalchemy.sql import select
s = cookies.select()
rp = connection.execute(s)
result = rp.fetchall()

for cookie in results:
    print(f'{cookie} \n')

(1, 'chocolate chip', 'http://some.aweso.me/cookie/recipe.html', 'CCO1', 12,
Decimal('0.50'))

(2, 'dark chocolate chip', 'https://some.aweso.me/cookie/recipe_dark.html',
'CCO2', 1, Decimal('0.75'))

(3, 'peanut butter', 'http://some.aweso.me/cookie/peanut.html', 'PBO1', 24,
Decimal('0.25'))

(4, 'oatmeal raisin', 'http://some.okay.me/cookie/raisin.html', 'EWW01', 100,
Decimal('1.00'))
```

0.3.3 ResultProxy

- a wrapper around a DBAPI cursor object and the goal is to make it easier to use and manipulate the results of a statement
- it could make handling query results easier by allowing access using an index, name or Column object
- all of the following statmeents result in the same thing: 'choclate chip'
- they each refrence the exact same data element in the first record of our results variable

```
[14]: # get the first row of the resultproxy
    first_row = results[0]

# access column by index
    first_row[1]

# access column by name
    first_row.cookie_name
```

```
# access column by Column object
first_row[cookies.c.cookie_name]
```

[14]: 'chocolate chip'

• we can also leverage the ResultProxy as an iterable and perform an action on each record returned without creatingt another variable to hold the results

```
[15]: # print the name of each cookie in our database

rp = connection.execute(s)

for record in rp:
    print(record.cookie_name)
```

chocolate chip dark chocolate chip peanut butter oatmeal raisin

- you can use the following methods to fetch results
- first()
 - returns the first record if there is one and closes the connection
- fetchone()
 - returns one row and leaves the cursor open for you to make additional fetch calls
- scalar()
 - returns a signle value if a query results in a single record with one column
- if you want to see the columns that are available in a result set you can use the keys() method to get a list of the column names

Tips for Good Production Code

- use the first method for getting a single record over both the fetchone and scalar methods
- use the iterable version of the ResultProxy over the fetchall and fetchone methods. It is more memory efficient and we tend to operate on the data one record at a time
- avoid the fetchone method, as it leaves connections open if you are not careful
- use the scalar method sparingly, as it raises errors if a query every returns more than one row and one column, which often gets missed during testing

0.3.4 Controlling The Columns in The Query

• to limit the fields that are returned from a query, we need to pass the columns we want into the select() method constructor as a list

```
[16]: s = select([cookies.c.cookie_name, cookies.c.quantity])
rp = connection.execute(s)

# returns the list of columns which is [u'cookie_name', u'quantity']
print('getting the list of columns')
```

```
print(rp.keys())
print('')
# returns only the first field
print('getting the first field')
result = rp.first()
print(result)
```

```
getting the list of columns
['cookie_name', 'quantity']
getting the first field
('chocolate chip', 12)
```

0.3.5 Ordering

- if we want the list to be returned in a particular order, we can chain an order_by() statement to our select
- desc() can be used to reverse the order

```
[17]: s = select([cookies.c.cookie_name, cookies.c.quantity])
s = s.order_by(cookies.c.quantity)
rp = connection.execute(s)

for cookie in rp:
    print(f'{cookie.quantity} - {cookie.cookie_name}')
1 - dark chocolate chip
```

```
1 - dark chocolate chip
12 - chocolate chip
24 - peanut butter
100 - oatmeal raisin
```

0.3.6 Limiting

• we have to use limit() method to limit

```
[18]: s = select([cookies.c.cookie_name, cookies.c.quantity])
s = s.order_by(cookies.c.quantity)
s = s.limit(2)
rp = connection.execute(s)

print([result.cookie_name for result in rp])
```

['dark chocolate chip', 'chocolate chip']

0.3.7 Built-In SQL Functions and Labels

• two of the most common database functions are SUM() and COUNT()

```
[19]: from sqlalchemy.sql import func
```

```
s = select([func.sum(cookies.c.quantity)])
rp = connection.execute(s)
# scalar is returning the left most column in the first record
print("returning the left most colum in the first record")
print(rp.scalar())
print("")
s2 = select([func.count(cookies.c.cookie name)])
rp = connection.execute(s2)
record = rp.first()
# this will show us the columns in the ResultProxy
print('these are the columns in the ResultProxy')
print(record.keys())
print("")
# the column name is autogenerated and commonly `<func_name>_<position>`
print("this is the record count")
print(record.count_1)
print("")
returning the left most colum in the first record
```

these are the columns in the ResultProxy
['count_1']

this is the record count

- label() function helps us label better so that the column name is not annoying and cumbersome
- you dont really want to be using count_1

```
[20]: s = select([func.count(cookies.c.cookie_name).label('inventory_count')])
    rp = connection.execute(s)
    record = rp.first()

print('these are the columns in the ResultProxy')
    print(record.keys())
    print("")

print(print("this is the record count"))
    print(record.inventory_count)
```

these are the columns in the ResultProxy

```
['inventory_count']
     this is the record count
     None
     4
     0.3.8 Filtering
        • filtering queries is done by adding a where() statement just like in SQL
        • we can chain multiple where() clauses
[21]: s = select([cookies]).where(cookies.c.cookie_name == 'chocolate chip')
      rp = connection.execute(s)
      record = rp.first()
      print('a list ofcolumns and values \n')
      for record in record.items():
          print(record)
          print('')
     a list of columns and values
     ('cookie_id', 1)
     ('cookie_name', 'chocolate chip')
     ('cookie_recipe_url', 'http://some.aweso.me/cookie/recipe.html')
     ('cookie_sku', 'CC01')
     ('quantity', 12)
     ('unit_cost', Decimal('0.50'))
[22]: # we can also find statements that contain particular words in them
      s = select([cookies]).where(cookies.c.cookie_name.like('%chocolate%'))
      rp = connection.execute(s)
      print('here are cookies types that begin with %chocolate% \n')
      for record in rp.fetchall():
          print(record.cookie_name)
     here are cookies types that begin with %chocolate%
     chocolate chip
```

dark chocolate chip

0.3.9 Clause Elements

- ClauseElements are just an entity we use in a clause, and they are typically columns in a table
- unlike columns, ClauseElements come with many additional capabilities
- like method is an example of this
- there are also negation versions of those methods called isnot()
- instead of using these, we can just use Operators

```
[23]: ![](images/02.png)
```

'[]' is not recognized as an internal or external command, operable program or batch file.

0.3.10 Operators

- you can use most operators
- you can also use /+, -, *, /, % for sting contraction

```
[24]: s = select([cookies.c.cookie_name, 'SKU-' + cookies.c.cookie_sku])

print('showing the cookies along with their SKU \n')
for row in connection.execute(s):
    print(row)
```

showing the cookies along with their SKU

```
('chocolate chip', 'SKU-CC01')
('dark chocolate chip', 'SKU-CC02')
('peanut butter', 'SKU-PB01')
('oatmeal raisin', 'SKU-EWW01')
```

• we can also use operators to compute values

showing inventory value by cookie

```
chocolate chip - 6.00 dark chocolate chip - 0.75
```

```
peanut butter - 6.00
oatmeal raisin - 100.00
```

0.3.11 Boolean Operators

- be careful because & binds more closely than <
- you want to use Conjuction instead of using & or < because it is more expressive

0.3.12 Conjuction

- when it is possible to chain multiple where() clauses together, it's often more readable and functional to use conjuctions to accomplish the desired effect
- conjuctions are and_(), or_(), not_()

using and() conjuction to show filtered data

peanut butter

0.4 Updating Data

- update statements can be created by either the update() function or the update() method on the table being updated
- if you dont have a where clause, you can update everything

```
[27]: from sqlalchemy import update

u = update(cookies).where(cookies.c.cookie_name == 'chocolate chip')
u = u.values(quantity=(cookies.c.quantity + 120))
result = connection.execute(u)

print('showing how many rows were updated')
print(result.rowcount, "\n")

s = select([cookies]).where(cookies.c.cookie_name == 'chocolate chip')
result = connection.execute(s).first()
```

```
print('showing the new updated results')
      for key in result.keys():
          print(f'{key}: {result[key]}')
     showing how many rows were updated
     showing the new updated results
     cookie_id: 1
     cookie_name: chocolate chip
     cookie_recipe_url: http://some.aweso.me/cookie/recipe.html
     cookie_sku: CC01
     quantity: 132
     unit_cost: 0.50
     0.4.1 Deleting Data
        • we can either use the delete() function or method on the table
        • unlike insert or update, delete takes no values, only optional where clause
[28]: from sqlalchemy import delete
      u = delete(cookies).where(cookies.c.cookie name == "dark chocolate chip")
      result = connection.execute(u)
      print('how many results have been found')
      print(result.rowcount)
      print("")
      s = select([cookies]).where(cookies.c.cookie_name == "dark chocolate chip")
      result = connection.execute(s).fetchall()
      print('fetching all instances of dark choclate chip')
      print(len(result))
     how many results have been found
     fetching all instances of dark choclate chip
[29]: # we are loading in some customer data
```

customer_list = [

{

```
'username': 'cookiemon',
    'email_address': 'mon@cookie.com',
    'phone': '111-111-1111',
    'password': 'password'
    },
    'username': 'cakeeater',
    'email_address': 'cakeeater@cake.com',
    'phone': '222-222-222',
    'password': 'password'
    },
    'username': 'pieguy',
    'email_address': 'guy@pie.com',
    'phone': '333-333-3333',
    'password': 'password'
    }
]
ins = users.insert()
result = connection.execute(ins, customer_list)
```

```
[30]: from sqlalchemy import insert
      # enter customer data
      ins = insert(orders).values(user_id=1, order_id=1)
      result = connection.execute(ins)
      ins = insert(line_items)
      order_items = [
          {
          'order_id': 1,
          'cookie_id': 1,
          'quantity': 2,
          'extended_cost': 1.00
          },
          'order_id': 1,
          'cookie_id': 3,
          'quantity': 12,
          'extended_cost': 3.00
      ]
      result = connection.execute(ins, order_items)
      ins = insert(orders).values(user_id=2, order_id=2)
      result = connection.execute(ins)
      ins = insert(line_items)
```

0.4.2 Joins

• select from() allows us to replace the entire from clause with a more specific one

```
(1, 'cookiemon', '111-111-1111', 'chocolate chip', 2, Decimal('1.00')) (1, 'cookiemon', '111-111-1111', 'peanut butter', 12, Decimal('3.00'))
```

• the sqlalchemy statement above looks like this

SELECT orders.order_id, users.username, users.phone, cookies.cookie_name, line_items.quantity, line_items.extended_cost FROM users JOIN orders ON users.user_id = orders.user_id JOIN line_items ON orders.order_id = line_items.order_id JOIN cookies ON cookies.cookie_id = line_items.cookie_id WHERE users.username = :username_1

- we can also use outerjoin to select from multiple tables
- sqlalchemy knows how to join the users and order tables because of the forgein key defined in the orders table

```
[32]: columns = [users.c.username, func.count(orders.c.order_id)]
      all_orders = select(columns)
      all_orders = all_orders.select_from(users.outerjoin(orders))
      all_orders = all_orders.group_by(users.c.username)
      result = connection.execute(all_orders).fetchall()
      for row in result:
          print(row)
     ('cakeeater', 1)
     ('cookiemon', 1)
     ('pieguy', 0)
     0.4.3 Aliases
        • SQLAlchemy allows the use of aliasing with the use of alias function or methods
     employee_table = Table(
          'employee', metadata,
         Column('id', Integer, primary_key=True),
         Column('manager', None, ForeignKey('employee.id')),
         Column('name', String(255)))
     111
     SELECT employee.name
     FROM employee, employee AS manager
     WHERE employee.manager_id = manager.id
         AND manager.name = 'Fred'
     111
     manager = employee_table.alias('mgr')
     stmt = select([employee_table.c.name],
             and_(employee_table.c.manager_id==manager.c.id,
                 manager.c.name=='Fred'))
     print(smtmt)
     manager = employee_table.alias()
     stmt = select([employee_table.c.name],
         and_(employee_table.c.manager_id==manager.c.id,
         manager.c.name=='Fred'))
         print(stmt)
     SELECT employee.name
     FROM employee, employee AS employee_1
     WHERE employee.manager_id = employee_1.id AND employee_1.name = ?
```

0.4.4 Grouping

• when grouping, you need one or more columns to group on and one or more columns that it makes sense to aggrate with count, sum, etc

```
[33]: columns = [users.c.username, func.count(orders.c.order_id)]
    all_orders = select(columns)
    all_orders = all_orders.select_from(users.outerjoin(orders))
    all_orders = all_orders.group_by(users.c.username)
    result = connection.execute(all_orders).fetchall()

    for row in result:
        print(row)

        ('cakeeater', 1)
        ('cookiemon', 1)
        ('pieguy', 0)

        0.4.5 Chainig

[34]: def get_orders_by_customer(cust_name):
        columns = [orders.c.order_id, users.c.username, users.c.phone,
```

cust_orders = cust_orders.select_from(joins)

```
cust_orders = cust_orders.where(users.c.username == cust_name)
  if shipped is not None:
        cust_orders = cust_orders.where(orders.c.shipped == shipped)
  result = connection.execute(cust_orders).fetchall()
  return result

get_orders_by_customer('cakeeater')
get_orders_by_customer('cakeeater', details=True)
get_orders_by_customer('cakeeater', shipped=True)
get_orders_by_customer('cakeeater', shipped=False)
get_orders_by_customer('cakeeater', shipped=False, details=True)
'''
```

0.5 Raw Queries

```
[35]: result = connection.execute("select * from orders").fetchall()
    print(result)

[(1, 1), (2, 2)]

[36]: from sqlalchemy import text

    stmt = select([users]).where(text("username='cookiemon'"))
    print(connection.execute(stmt).fetchall())

[(1, None, 'cookiemon', 'mon@cookie.com', '111-111-1111', 'password',
    datetime.datetime(2021, 2, 22, 8, 38, 51, 536158), datetime.datetime(2021, 2, 22, 8, 38, 51, 536158))]
```

0.6 Exceptions and Transactions

- the most common errors are ${\tt AttributeErrors}$ and ${\tt IntegrityErrors}$

```
CheckConstraint('quantity > 0', name='quantity_positive')
)
users = Table('users', metadata,
    Column('user_id', Integer(), primary_key=True),
    Column('username', String(15), nullable=False, unique=True),
    Column('email_address', String(255), nullable=False),
    Column('phone', String(20), nullable=False),
    Column('password', String(25), nullable=False),
    Column('created_on', DateTime(), default=datetime.now),
    Column('updated_on', DateTime(), default=datetime.now,
            onupdate=datetime.now)
)
orders = Table('orders', metadata,
    Column('order_id', Integer()),
    Column('user_id', ForeignKey('users.user_id')),
    Column('shipped', Boolean(), default=False)
)
line_items = Table('line_items', metadata,
    Column('line_items_id', Integer(), primary_key=True),
    Column('order_id', ForeignKey('orders.order_id')),
    Column('cookie id', ForeignKey('cookies.cookie id')),
    Column('quantity', Integer()),
    Column('extended cost', Numeric(12, 2))
engine = create_engine('sqlite:///:memory:')
metadata.create_all(engine)
connection = engine.connect()
```

0.6.1 AttributeError

occurs when you attempt to access a column on a ResultProxy that isn't present

```
[5]: # this is an example of an attributeError

from sqlalchemy import select, insert

ins = insert(users).values(
    username="cookiemon",
    email_address="mon@cookie.com",
    phone="111-111-1111",
    password="password"
)
```

```
result = connection.execute(ins)
s = select([users.c.username])
results = connection.execute(s)

for result in results:
    print(result.username)
    print(result.password)
```

```
Traceback (most recent call last)
IntegrityError
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
→py in _execute_context(self, dialect, constructor, statement, parameters,
→*args)
   1275
                        if not evt_handled:
-> 1276
                            self.dialect.do_execute(
   1277
                                cursor, statement, parameters, context
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
→py in do_execute(self, cursor, statement, parameters, context)
            def do_execute(self, cursor, statement, parameters, context=None):
--> 608
                cursor.execute(statement, parameters)
    609
IntegrityError: UNIQUE constraint failed: users.username
The above exception was the direct cause of the following exception:
IntegrityError
                                          Traceback (most recent call last)
<ipython-input-5-02115079d03f> in <module>
     10 )
     11
---> 12 result = connection.execute(ins)
     13 s = select([users.c.username])
     14 results = connection.execute(s)
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
→py in execute(self, object_, *multiparams, **params)
   1009
   1010
                else:
-> 1011
                   return meth(self, multiparams, params)
   1012
   1013
            def _execute_function(self, func, multiparams, params):
```

```
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
 →py in _execute_on_connection(self, connection, multiparams, params)
            def execute on connection(self, connection, multiparams, params):
    297
                if self.supports_execution:
--> 298
                    return connection._execute_clauseelement(self, multiparams,
 →params)
    299
                else:
    300
                    raise exc.ObjectNotExecutableError(self)
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
→py in execute clauseelement(self, elem, multiparams, params)
   1122
   1123
-> 1124
                ret = self._execute_context(
   1125
                    dialect.
   1126
                    dialect.execution_ctx_cls._init_compiled,
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
→py in _execute_context(self, dialect, constructor, statement, parameters, ____
→*args)
   1314
   1315
                except BaseException as e:
                    self._handle_dbapi_exception(
-> 1316
   1317
                        e, statement, parameters, cursor, context
   1318
                    )
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
→py in _handle_dbapi_exception(self, e, statement, parameters, cursor, context
                        util.raise (newraise, with traceback=exc info[2],
   1508
 \rightarrowfrom =e)
   1509
                    elif should wrap:
-> 1510
                        util.raise (
   1511
                            sqlalchemy_exception, with_traceback=exc_info[2],_
 \hookrightarrowfrom_=e
   1512
                        )
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s |
 →py in raise_(***failed resolving arguments***)
    180
    181
                try:
--> 182
                    raise exception
    183
                finally:
    184
                    # credit to
```

```
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s | lalchemy\eng
→py in _execute_context(self, dialect, constructor, statement, parameters,
→*args)
  1274
                                    break
   1275
                        if not evt_handled:
-> 1276
                            self.dialect.do_execute(
   1277
                                cursor, statement, parameters, context
   1278
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\eng
→py in do_execute(self, cursor, statement, parameters, context)
    606
    607
            def do_execute(self, cursor, statement, parameters, context=None):
--> 608
                cursor.execute(statement, parameters)
    609
    610
            def do_execute_no_params(self, cursor, statement, context=None):
IntegrityError: (sqlite3.IntegrityError) UNIQUE constraint failed: users.userna e
[SQL: INSERT INTO users (username, email address, phone, password, created on,
→updated_on) VALUES (?, ?, ?, ?, ?, ?)]
[parameters: ('cookiemon', 'mon@cookie.com', '111-111-1111', 'password', |
→ '2021-02-22 09:31:24.772878', '2021-02-22 09:31:24.772878')]
(Background on this error at: http://sqlalche.me/e/13/gkpj)
```

0.6.2 IntegrityError

- occurs wheen we try to do something that would violate the constraints configured on a Column or Table
- this type of error is commonly encountered in cases where you require something to be unique -for example, if you attempt to create two users with the same username, an IntegrityError will be thrown because usernames in our users table must be unique

```
[9]: s = select([users.c.username])
    connection.execute(s).fetchall()

    [(u'cookiemon',)]

    ins = insert(users).values(
        username="cookiemon",
        email_address="damon@cookie.com",
        phone="111-111-1111",
        password="password"
    )
```

```
IntegrityError
                                          Traceback (most recent call last)
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
→py in _execute_context(self, dialect, constructor, statement, parameters,
→*args)
                        if not evt handled:
   1275
-> 1276
                            self.dialect.do_execute(
   1277
                                cursor, statement, parameters, context
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
 →py in do_execute(self, cursor, statement, parameters, context)
            def do_execute(self, cursor, statement, parameters, context=None):
--> 608
                cursor.execute(statement, parameters)
    609
IntegrityError: UNIQUE constraint failed: users.username
The above exception was the direct cause of the following exception:
IntegrityError
                                          Traceback (most recent call last)
<ipython-input-9-db414867bc94> in <module>
     11 )
     12
---> 13 result = connection.execute(ins)
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
→py in execute(self, object_, *multiparams, **params)
   1009
   1010
                else:
-> 1011
                    return meth(self, multiparams, params)
   1012
   1013
            def _execute_function(self, func, multiparams, params):
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\sql
 →py in _execute_on_connection(self, connection, multiparams, params)
    296
            def _execute_on_connection(self, connection, multiparams, params):
    297
                if self.supports execution:
                    return connection._execute_clauseelement(self, multiparams,
--> 298
 →params)
    299
                else:
    300
                    raise exc.ObjectNotExecutableError(self)
```

```
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
 →py in _execute_clauseelement(self, elem, multiparams, params)
   1122
   1123
-> 1124
                ret = self._execute_context(
   1125
                    dialect.
   1126
                    dialect.execution_ctx_cls._init_compiled,
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\eng
 →py in _execute_context(self, dialect, constructor, statement, parameters,
→*args)
   1314
   1315
                except BaseException as e:
-> 1316
                    self._handle_dbapi_exception(
   1317
                        e, statement, parameters, cursor, context
   1318
                    )
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
 →py in _handle_dbapi_exception(self, e, statement, parameters, cursor, context
                        util.raise (newraise, with traceback=exc info[2],
   1508
 →from =e)
   1509
                    elif should_wrap:
-> 1510
                        util.raise_(
                            sqlalchemy_exception, with_traceback=exc_info[2],_
   1511
 →from =e
   1512
                        )
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\uti
→py in raise_(***failed resolving arguments***)
    180
    181
                try:
--> 182
                    raise exception
    183
                finally:
    184
                    # credit to
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\state{lalchemy\eng}
→py in _execute_context(self, dialect, constructor, statement, parameters,
 →*args)
   1274
                                    break
                        if not evt handled:
   1275
-> 1276
                            self.dialect.do_execute(
                                cursor, statement, parameters, context
   1277
                            )
   1278
```

```
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s | lalchemy\eng
→py in do execute(self, cursor, statement, parameters, context)
    606
    607
            def do_execute(self, cursor, statement, parameters, context=None):
--> 608
                cursor.execute(statement, parameters)
    609
            def do_execute_no_params(self, cursor, statement, context=None):
    610
IntegrityError: (sqlite3.IntegrityError) UNIQUE constraint failed: users.userna e
[SQL: INSERT INTO users (username, email_address, phone, password, created_on,_
→updated_on) VALUES (?, ?, ?, ?, ?)]
[parameters: ('cookiemon', 'damon@cookie.com', '111-111-1111', 'password', __
→'2021-02-22 09:34:48.127762', '2021-02-22 09:34:48.127762')]
(Background on this error at: http://sqlalche.me/e/13/gkpj)
```

0.6.3 Handling Errors

• wrap as little code as possible in a try/except block

```
[12]: from sqlalchemy.exc import IntegrityError
   ins = insert(users).values(
        username="cookiemon",
        email_address="damon@cookie.com",
        phone="111-111-1111",
        password="password"
)

try:
    result = connection.execute(ins)
   except IntegrityError as error:
        print(error.orig.message, error.params)
```

```
607
            def do_execute(self, cursor, statement, parameters, context=None):
--> 608
                cursor.execute(statement, parameters)
    609
IntegrityError: UNIQUE constraint failed: users.username
The above exception was the direct cause of the following exception:
                                          Traceback (most recent call last)
IntegrityError
<ipython-input-12-b821eef8ec75> in <module>
      9 try:
---> 10
            result = connection.execute(ins)
     11 except IntegrityError as error:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
 →py in execute(self, object , *multiparams, **params)
                else:
-> 1011
                    return meth(self, multiparams, params)
   1012
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\sql
→py in _execute_on_connection(self, connection, multiparams, params)
    297
                if self.supports_execution:
--> 298
                    return connection._execute_clauseelement(self, multiparams,
→params)
    299
                else:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s
→py in _execute_clauseelement(self, elem, multiparams, params)
   1123
-> 1124
                ret = self._execute_context(
   1125
                    dialect.
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s|lalchemy\eng
 →py in execute context(self, dialect, constructor, statement, parameters,
→*args)
   1315
                except BaseException as e:
-> 1316
                    self._handle_dbapi_exception(
   1317
                        e, statement, parameters, cursor, context
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\silalchemy\eng
 →py in _handle_dbapi_exception(self, e, statement, parameters, cursor, context
   1509
                    elif should wrap:
```

```
-> 1510
                        util.raise_(
   1511
                            sqlalchemy_exception, with_traceback=exc_info[2],_
 →from_=e
c:
→\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\uti
→py in raise_(***failed resolving arguments***)
    181
                try:
--> 182
                    raise exception
    183
                finally:
c:
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\eng
→py in execute context(self, dialect, constructor, statement, parameters,
→*args)
   1275
                        if not evt_handled:
-> 1276
                            self.dialect.do_execute(
   1277
                                cursor, statement, parameters, context
 →\users\vicktree\appdata\local\programs\python\python38-32\lib\site-packages\s lalchemy\eng
→py in do execute(self, cursor, statement, parameters, context)
            def do_execute(self, cursor, statement, parameters, context=None):
    607
--> 608
                cursor.execute(statement, parameters)
    609
IntegrityError: (sqlite3.IntegrityError) UNIQUE constraint failed: users.usernale
[SQL: INSERT INTO users (username, email address, phone, password, created on,
→updated_on) VALUES (?, ?, ?, ?, ?)]
[parameters: ('cookiemon', 'damon@cookie.com', '111-111-1111', 'password',
\rightarrow '2021-02-22 09:38:06.519785', '2021-02-22 09:38:06.519785')]
(Background on this error at: http://sqlalche.me/e/13/gkpj)
During handling of the above exception, another exception occurred:
AttributeError
                                          Traceback (most recent call last)
<ipython-input-12-b821eef8ec75> in <module>
            result = connection.execute(ins)
     11 except IntegrityError as error:
            print(error.orig.message, error.params)
AttributeError: 'IntegrityError' object has no attribute 'message'
```

0.6.4 Transactions

- transactions are a way to ensure that multiple database statements succeed or fail as a group
- when we start a transaction, we record the current state of our databas

- then we can execute multiple SQL statements
- if all the SQL statements in the transaction succeed, the database continues on normally and we discard the prior database state

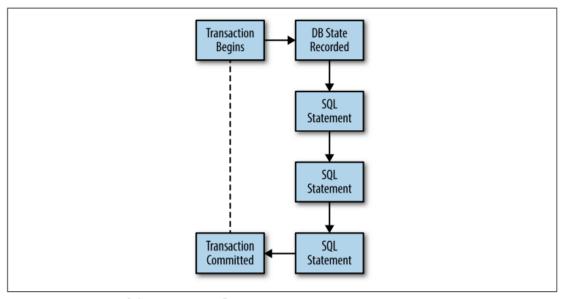


Figure 3-1. Successful transaction flow

- however if one or more of these statements fail,
- we can catch that error and use the prior state to roll back any statements that succeedded

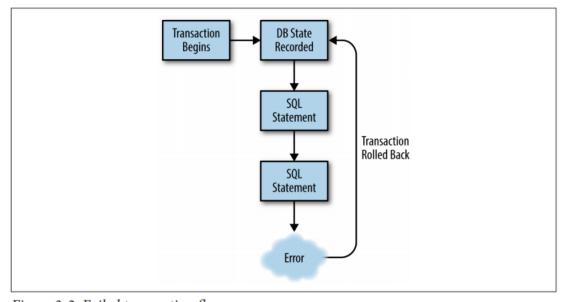


Figure 3-2. Failed transaction flow

```
[3]: # Setting up the Transaction enviorment

from datetime import datetime
from sqlalchemy import (MetaData, Table, Column, Integer, Numeric, String,
```

```
DateTime, ForeignKey, Boolean, create_engine,
                        CheckConstraint)
metadata = MetaData()
cookies = Table('cookies', metadata,
   Column('cookie_id', Integer(), primary_key=True),
   Column('cookie_name', String(50), index=True),
   Column('cookie recipe url', String(255)),
   Column('cookie_sku', String(55)),
   Column('quantity', Integer()),
   Column('unit_cost', Numeric(12, 2)),
   # Note the check we are doing here, before we actually send out shipment
   # cookies, we need to check that we have that quantity
   CheckConstraint('quantity >= 0', name='quantity_positive')
)
users = Table('users', metadata,
   Column('user_id', Integer(), primary_key=True),
   Column('username', String(15), nullable=False, unique=True),
   Column('email_address', String(255), nullable=False),
   Column('phone', String(20), nullable=False),
   Column('password', String(25), nullable=False),
   Column('created_on', DateTime(), default=datetime.now),
   Column('updated on', DateTime(), default=datetime.now,
           onupdate=datetime.now)
)
orders = Table('orders', metadata,
   Column('order_id', Integer()),
   Column('user_id', ForeignKey('users.user_id')),
   Column('shipped', Boolean(), default=False)
)
line_items = Table('line_items', metadata,
   Column('line_items_id', Integer(), primary_key=True),
   Column('order_id', ForeignKey('orders.order_id')),
   Column('cookie id', ForeignKey('cookies.cookie id')),
   Column('quantity', Integer()),
   Column('extended_cost', Numeric(12, 2))
)
engine = create_engine('sqlite:///:memory:')
metadata.create_all(engine)
connection = engine.connect()
```

```
from sqlalchemy import select, insert, update
ins = insert(users).values(
    username="cookiemon",
    email_address="mon@cookie.com",
    phone="111-111-1111",
    password="password"
)
result = connection.execute(ins)
ins = cookies.insert()
inventory_list = [
    {
        'cookie_name': 'chocolate chip',
        'cookie_recipe_url': 'http://some.aweso.me/cookie/recipe.html',
        'cookie_sku': 'CC01',
        'quantity': '12',
        'unit_cost': '0.50'
   },
        'cookie_name': 'dark chocolate chip',
        'cookie_recipe_url': 'http://some.aweso.me/cookie/recipe_dark.html',
        'cookie sku': 'CCO2',
        'quantity': '1',
        'unit cost': '0.75'
    }
]
result = connection.execute(ins, inventory_list)
ins = insert(orders).values(user_id=1, order_id='1')
result = connection.execute(ins)
ins = insert(line_items)
order_items = [
     {
         'order_id': 1,
         'cookie_id': 1,
         'quantity': 9,
         'extended_cost': 4.50
     }
result = connection.execute(ins, order_items)
ins = insert(orders).values(user_id=1, order_id='2')
result = connection.execute(ins)
ins = insert(line_items)
```

```
[4]: def ship_it(order_id):
    s = select([line_items.c.cookie_id, line_items.c.quantity])
    s = s.where(line_items.c.order_id == order_id)
    cookies_to_ship = connection.execute(s)
    for cookie in cookies_to_ship:
        u = update(cookies).where(cookies.c.cookie_id==cookie.cookie_id)
        u = u.values(quantity = cookies.c.quantity - cookie.quantity)
        result = connection.execute(u)
    u = update(orders).where(orders.c.order_id == order_id)
    u = u.values(shipped=True)
    result = connection.execute(u)
    print("Shipped order ID: {}".format(order_id))
```

```
[5]: # running ship_it on the first order
ship_it(1)
s = select([cookies.c.cookie_name, cookies.c.quantity])
connection.execute(s).fetchall()
```

Shipped order ID: 1

- [5]: [('chocolate chip', 3), ('dark chocolate chip', 1)]
 - if we tried to run the second order, we would get IntegrityError because we would not have enough chocolate chip cookies
 - the problem is however, it did change the database value for one of the cookies before throwing an exception
 - we need to use a transaction

```
[7]: from sqlalchemy.exc import IntegrityError

def ship_it(order_id):
    s = select([line_items.c.cookie_id, line_items.c.quantity])
```

```
s = s.where(line_items.c.order_id == order_id)
transaction = connection.begin()
cookies_to_ship = conenction.execute(s).fetchall()

try:
    for cookie in cookies_to_ship:
        u = update(cookies).where(cookies.c.cookie_id == cookie.cookie_id)
        u = u.values(quantity == cookies.c.quantity - cookie.quantity)
        result = connection.execute(u)
    u = update(orders).where(orders.c.order_id == order_id)
    u = u.values(shipped=True)
    print("Shipped order ID: {}".format(order_id))
    transaction.commit()
except IntegrityError as error:
    transaction.rollback()
    print(error)
```

- we import IntegrityError so we can handle its exception
- we start the transaction by using begin()
- we use commit() to commit the transactions if no error occurs
- we use rollback() incase an error occurs

0.7 Testing

0.7.1 Testing with a Database

- we will have an app.py for the application logic and an db.py file that contains our database tables and connections
- in the db.py file, we have our database set up via the DataAccessLayer class
- the class is used to initialize a database cehema and connect it to an engine
- the DataAccessLayer class is initialized without an engine and a connection in the dal variable

```
[8]: # app.py
     # this is our DataAccessLayer instance from the db.py file
     #from db import dal
     from sqlalchemy.sql import select
     def get_orders_by_customers(cust_name, shipped=None, details=False):
         columns = [dal.orders.c.order_id, dal.users.c.username, dal.users.c.phone]
         # because our tables are inside of the dal object, we access them from here
         joins = dal.users.join(dal.orders)
         if details:
             columns.extend([dal.cookies.c.cookie_name,
                            dal.line_items.c.quantity,
                            dal.line_items.c.extended_cost])
         joins = joins.joins(dal.line_items).join(dal.cookies)
         cust_orders = select(columns)
         cust_orders = cust_orders.select_from(joins).where(
         dal.users.c.username == cust name)
         if shipped is not None:
             cust_orders = cust_orders.where(dal.orders.c.shipped == shipped)
         return dal.connection.execute(cust_orders).fetchall()
```

- notice in the file above we have three inputs cust_name, shipped and details
 - cust_name can be blank, a string containing the name of a valid customer, or a string that does not contain the name of a valid customeer
 - shipped can be None, True, False
 - details can be True or False
- inorder to test all possible combinations, we need 12 (3 * 2 * 2) tests

```
[13]: import unittest

# unittest requires test classes inherited from unittest.TestCase
class TestApp(unittest.TestCase):

# setupClass method is run once for the entire test class
@classmethod
    def setUpClass(cls):

# a connection is initialized for an in-memory db
    dal._db_init('sqlite:///:memory:')

# this function fills database with dummy data
    # prep_db()
```

- below we are loading in data
- note that each method begins with a test_ in its name
- also note that we are using an assertEqual keyword

```
[12]: def test orders by customer blank(self):
          results = get_orders_by_customer('')
          self.assertEqual(results, [])
      def test_orders_by_customer_blank_shipped(self):
          results = get_orders_by_customer('', True)
          self.assertEqual(results, [])
      def test_orders_by_customer_blank_notshipped(self):
          results = get_orders_by_customer('', False)
          self.assertEqual(results, [])
      def test_orders_by_customer_blank_details(self):
          results = get_orders_by_customer('', details=True)
          self.assertEqual(results, [])
      def test_orders_by_customer_blank_shipped_details(self):
          results = get_orders_by_customer('', True, True)
          self.assertEqual(results, [])
      def test_orders_by_customer_blank_notshipped_details(self):
          results = get_orders_by_customer('', False, True)
```

```
self.assertEqual(results, [])
```

• run the tests above with python -m unittest test_app

0.7.2 Mocking

- Mocking is a powerful tool to use when you have a test environment where creating a test database doesnt make sense or simply isnt feasible
- it could be useful to mock out the SQLAlchemy code to return the value you want
- this allows you to focus on the surrounding logic

```
[14]: import sys
!{sys.executable} -m pip install mock
import mock
```

Collecting mock

Downloading mock-4.0.3-py3-none-any.whl (28 kB) Installing collected packages: mock Successfully installed mock-4.0.3

WARNING: You are using pip version 20.2.3; however, version 21.0.1 is available. You should consider upgrading via the

'c:\users\vicktree\appdata\local\programs\python\python39\python.exe -m pip
install --upgrade pip' command.

- Mock has a patch function that will let us replace a given object in a Python file with a MagicMock that we can control from our test
- a MagicMock is a special type of python object that tracks how it is used and allows us to define how it behaves based on how it is used
- below we are going to use the patch method as a decorator to replace the connection part of the dal object
- because the dal object is imported, we will need to patch it inside the app module
- this will get passed into test function as an argument
- now that we have mock object, we can set a return value for the execute method, which in this case should be nothing but a chained fetchall method whos return value is the data we want to test with

```
[17]: from unittest import mock
  from decimal import Decimal
  import mock

class TestApp(unittest.TestCase):
    cookie_orders = [(u'wlk001', u'cookiemon', u'111-111-1111')]
    cookie_details = [
        (u'wlk001', u'cookiemon', u'111-111-1111',
        u'dark chocolate chip', 2, Decimal('1.00')),
        (u'wlk001', u'cookiemon', u'111-111-1111',
        u'oatmeal raisin', 12, Decimal('3.00'))
```

```
patching dal.connection in the app module with a mock
that patch is passed into the test function as `mock_conn`

"""

Gmock.patch('app.dal.connection')
def test_orders_by_customer(self, mock_conn):

"""

we set the return value of the execute method to the chained returne
value of the fetchall method, which we set to self.cookie_order

"""

mock_conn.execute.return_value.fetchall.return_value = self.cookie_orders

"""

now we call the test function where the `dal. conenction will be mocked
and return the value we set in the prior step

"""

result = get_orders_by_customer('cookiemon')
self.assertEqual(result, self.cookie_orders)
```

- below we will use the mock.patch decorator and mock out the select object
- we will have to mock out all the chained query element return values, which in this query are the select_from and where clauses

0.8 Reflections

- Reflection allows us to handle an existing database with SQLAlchemy without the need to re-create all the schema in Python
- Reflection is a technique that allows us to populate a SQLAlchemy object from an existing database
- instead of defining columns by hand, we are going to use the autoload and autoload_with keyword argument this will reflect
- this will reflect the schema information into the metadata object and store a refrence to the table in the artist variable

you have to be careful about reflecting because what can happen is that if you reflect one
table and it relies on another table which has not been reflected, then SQLAlchemy will drop
that

Note, the code below will not work because you need the chinook database in the sqlite directory

0.8.1 Reflecting Individual Tables

```
[25]: def reflect_table(self):
          from sqlalchemy import MetaData, create_engine
          metadata = MetaData()
          # you need to have the chinook database which comes with a bunch of data
          engine = create_engine('sqlite://Chinook_Sqlite.sqlite')
          artist = Table('Artist', metadata, autoload=True, autoload_with=engine)
          # this code fetches 10 artists
          artist.columns.keys()
          from sqlalchemy import select
          s = Select([artist]).limit(10)
          engine.execute(s).fetchall()
          album = Table('Album', metadata, autoload=True, autoload with=engine)
          # this lets you view the meta data
          metadata.tables['album']
          above statement prints:
          Table('album',
           MetaData(bind=None),
           Column('AlbumId', INTEGER(), table=<album>, primary key=True,,,
       \rightarrow nullable=False),
           Column('Title', NVARCHAR(length=160), table=<album>, nullable=False),
```

```
Column('ArtistId', INTEGER(), table=<album>, nullable=False),
 schema=None)
111
# check if reflection occured
album.forein_keys
# getting back dropped relationships
from sqlalchemy import ForeignKeyConstrain
album.append_constraint(
    ForeignKeyConstraint(['ArtistId'], ['artist.ArtistId'])
)
# this lets you view the meta data
metadata.tables['album']
# we can see database relations by using the `str` keyword
str(artist.join(album))
111
'artist JOIN album ON artist."ArtistId" = album."ArtistId"'
```

0.8.2 Reflecting a Whole Database

- inorder to reflect a whole database, we can use the reflect method on the metadata object
- the reflect method will scan everything available on the engine supplied and reflect everything it can

0.8.3 Query Building with Reflected Objects

- we need a way to refer to the test of the tables that were reflected when we reflected the entire database
- we will need a way to refer to them in our query
- we can do this by assigning them to a variable from the tables attribute of the metadata

```
[28]: def build_reflected_query():
    # establish a variable to be refrence to the table
    playlist = metadata.tables['Playlist']
    from sqlalchemy import select

# use the variable in the query
    s = select([playlist]).limit(10)
    engine.execute(s).fetchall()
```

• by assigning the reflected tables we want to use into a variable, we can use them as we normall used